

On the Design of Autonomic, Decentralized VPNs

David Isaac Wolinsky, Kyungyong Lee, P. Oscar Boykin, Renato Figueiredo
University of Florida

Abstract—Decentralized and P2P (peer-to-peer) VPNs (virtual private networks) are becoming quite popular to connect users in small to medium collaborative environments, such as academia, businesses, and homes. In the realm of VPNs, there exist centralized, decentralized, and P2P solutions. Centralized systems require that at least a single entity providing an infrastructure for everyone, decentralized approaches allow more than one entity to share the responsibility for this infrastructure, while existing P2P approaches rely on a centralized infrastructure but allow users to bypass the centralized infrastructure to form direct low latency, high throughput links. In this paper, we describe a novel VPN architecture that can claim to be both decentralized and P2P using methods that limit the entry barrier for VPN deployment. Our solution extends existing work, IPOP (IP over P2P), to address challenges of configuration, management, and bootstrapping to create cryptographically secure, self-contained P2P systems. In doing so, we present the first implementation and analysis of a P2P system secured by DTLS (datagram transport layer security) along with decentralized techniques for revoking user access.

I. INTRODUCTION

A Virtual Private Network (VPN) provides the illusion of a local area network (LAN) spanning a wide area network (WAN) by creating secure¹ communication links amongst participants. Common uses of VPNs include secure access to enterprise network resources from remote/insecure locations, connecting distributed resources from multiple sites, and establishing virtual LANs for multiplayer video games and media sharing over the Internet.

The architecture described in this paper addresses usage scenarios where VPNs are desired but complexity in deployment and management limits their applicability. Such as collaborative academic environments linking individuals spanning multiple institutions, where coordinated configuration of network infrastructure across the different sites is often impractical. Another example is the small/medium business (SMB) environment, where it is often desirable to securely connect desktops and servers across distributed sites without incurring the complexity or management costs of traditional VPNs. Such a VPN could be used to enable extended families to share media, such as family videos and pictures, with each other, where existing VPNs may be too complicated or the hosting a centralized service may be undesirable.

Explicitly, the model considered in this paper is motivated from our Archer [1] project. Archer provides a dynamic and decentralized grid environment for computer architecture researchers to share and access voluntary compute cycles with each other. Use of centralized systems would limit the scope

¹For the remainder of this paper, unless explicitly stated otherwise, security implies encryption and mutual authentication between peers.

of Archer and require dedicated administration from multiple parties, whereas decentralized solutions require manual configuration of links between peers, which is beyond the scope of the target users. Current P2P VPN approaches either lack scalability or proper security components to be useful for VPN approaches.

We began our original foray into user-friendly VPN approaches with IPOP [2]. Previous work on IPOP focused on the routing mechanisms and address allocation with multiple virtual networks (VNs) sharing a single P2P overlay. Deploying a private overlay would require the user to deploy their own infrastructure including necessary security underlay, such as IPsec, at which point, at which point a P2P VPN becomes irrelevant. Sharing an overlay has significant drawbacks. A misconfigured peer could easily disable the entire overlay, rendering all VPNs useless, and the system would have to be recreated as there exists no methods to remove the peer from such a system. Without a shared overlay, each VPN would need to deploy a P2P infrastructure prior to the VPN system, at which point, users may reconsider the approach and prefer a more traditional centralized approach.

To make an easy-to-use, fully decentralized and P2P VPN, we extend the IPOP concept to support bootstrapping from public infrastructures and overlays into private P2P overlays whose membership is limited to an individual VPN user base. The key concept behind our work is based upon Castro et al. [3], who suggests a single overlay for bootstrapping service overlays. A private overlay by itself is secured only by obfuscation, malicious peers may eventually discover, join, and violate the overlay. As such, we present the first implementation and evaluation of an overlay with secure communication between end points in the P2P overlay and communication inside the P2P overlay. Security without revocation for misbehaving peers can ultimately lead to corrupted overlays but current revocation techniques rely on centralized systems, we present means to revoke using the P2P overlay itself. We call the completed system and the interface used to administrate it GroupVPN, a novel VPN architecture that extends IPOP to create a novel secured decentralized P2P VPN.

The rest of this paper is organized as follows. Section II describes the current IPOP architecture. Current IPOP systems lack security, we present a framework for securing both IPOP and the P2P overlay in Section V. Section VI overviews our approaches for decentralized certificate revocation. The complete system, GroupVPN, is presented in Section VII. Section VIII compares and contrasts our work with related work. Section IX concludes the paper.

II. BACKGROUND

This section describes the basic construction of IPOP, a structured P2P virtual network including background on structured overlays, address allocation and discovery, and connectivity.

A. P2P Overlays

The type of P2P overlay chosen for a VPN will have an effect on how easy the VPN is to program, deploy, secure, how efficient it is, and how it will deal with growth. The two primary infrastructures for P2P overlays are unstructured and structured systems. Unstructured systems use mechanisms such as global knowledge, broadcasts, or stochastic techniques [4] to search the overlay, as the system grows, maintaining this state and searching for things requires complex algorithms and mechanisms to retain efficiency. Alternatively, structured approaches maintain provide guaranteed search times typically with a lower bound of $O(\log N)$ regardless of network size. In terms of complexity, for small systems, unstructured systems may be easier to implement but as the system grows it may become inefficient.

IPOP uses a structured P2P framework named Brunet [5], which is based upon Symphony [6]. Structured systems are able to make guarantees about efficiencies by self-organizing into well-defined topologies, such as a 1 dimensional ring or a hypercube, each member having a unique identifier, a P2P address. Links in the overlay are made to guarantee the efficient search time. In addition, Brunet automatically creates links between heavily trafficked peers for efficiency in IPOP.

A key component of most structured overlays is support for decentralized storage/retrieval of information called a distributed hash table (DHT). The DHT builds upon the existence of a P2P address space. All peers in a structured system have a unique, uniformly distributed P2P address. A DHT maps look up values or keys usually by a hashing function into the P2P address space. While there are various forms of fault tolerance, a minimalist DHT stores values at the node whose address is closest to the value's key. DHTs can be used coordinate organization and discovery of resources, making them attractive for self-configuration and organization in decentralized collaborative environments. As explained in the next section, IPOP, uses a DHT to coordinate decentralized organization.

B. Connecting to the VPN

To connect to IPOP, a peer need only connect to an existing Brunet infrastructure. Many IPOP systems can coexist sharing a single overlay. The motivation for doing so is that bootstrapping a P2P system can be challenging, requiring users to understand concepts such as IP addresses and ports, as well as having access to a public IP address or being able to configure a router or firewall to enable inbound connections.

A peer connected to IPOP's P2P infrastructure can take advantage of its support for NAT traversal through hole punching [7]. When performing hole punching, peers first obtain mappings of their private IP address and port to their

public IP address and port and then exchange them over a shared medium, in this case the P2P overlay. The peers attempt to simultaneously form connections with each other, tricking NATs and firewalls into allowing inbound connections, because the NAT believed an outbound flow already exists. In case peers cannot establish direct connectivity, messages can be relayed through the P2P overlay to each other though with added latency and limited throughput.

This approach not only enables peers behind NATs and firewalls to seamlessly connect to each other, without requiring peers to host their own bootstrap servers. If a peer were to host their own bootstrap servers, they first need a public IP address and bind the application to a port on that system. At which point, they could share the IP, port pair with other peers in the VPN. Though if they were to go offline or their IP address were to change the P2P infrastructure, new users would be unable to join.

C. Organization

In the context of VPNs, structured overlays can handle organization of the network space, address allocation and discovery, decentrally through the use of a DHT, such systems have been proposed in [8], [9]. Membership in the VPN includes a matching membership in the structured overlay, thus all VPN peers have a P2P address. To address the challenges of having multiple VPNs in the same overlay, each IPOP group has its own namespace, reducing the likelihood of overlap. To enable scalable and decentralized address allocation and discovery, peers store mappings of IP address to P2P address into the DHT, typically of the form $\text{hash}(\text{namespace} + \text{IP}) = \text{P2Paddress}$. Thus a peer attempting to allocate an address will insert this key, value pair into the overlay. The first peer to do this will be the owner of the IP address allocation. Therefore the DHT must support atomic writes into the DHT.

Mechanisms to self-configure the IP address and network parameters of the local system can be provided by DHCP (decentralized host configuration protocol), manually configuring the IP address, or the VPN hooking into OS APIs. Address discovery is initiated when an outgoing packet for a remote peer arrives at the VPN software. At which point, the VPN will query the DHT with the IP address to obtain the owner's P2P address and forward the packet to the destination. Discussion on both these topics is further covered in our previous work [10].

III. EVALUATION TOOLS

As mentioned earlier, IPOP uses Brunet as the underlying P2P infrastructure for connectivity. Brunet has been in active development for the past 5 years. Prior to releasing updates, the code is run on PlanetLab [11] for a period of 2 weeks. PlanetLab consists of nearly 1,000 resources distributed across Earth. In practical application, though, roughly 40% of the resources are unavailable at any given time and the remaining behave somewhat unpredictably. This unpredictability provides behavior, that many times, is far more effective at finding bugs than users tests.

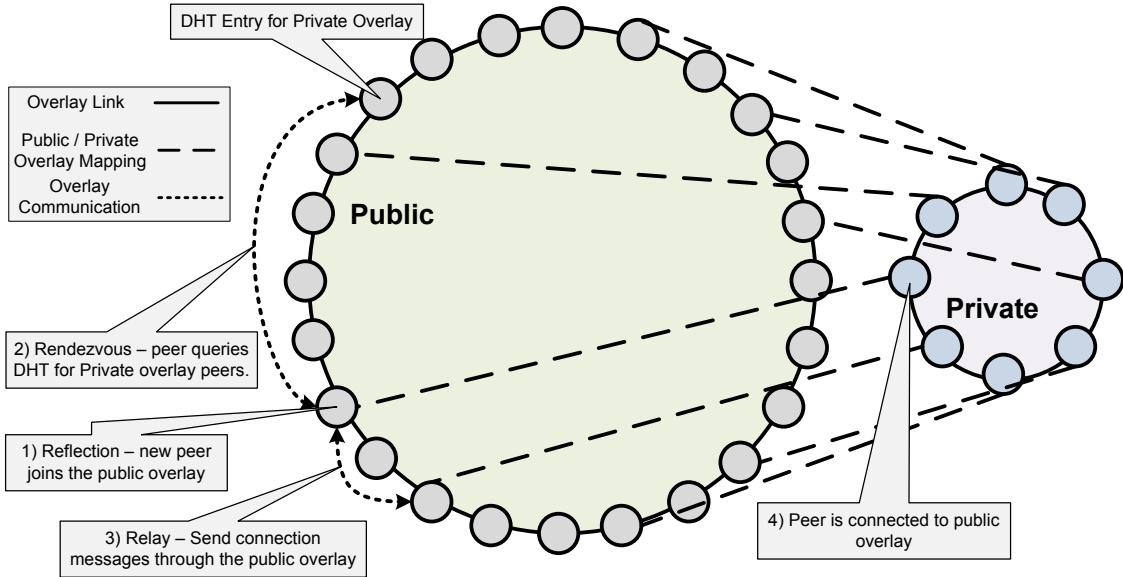


Fig. 1. Bootstrapping a private overlay using Brunet

Deploying on PlanetLab takes approximately 15 minutes to bootstrap all the resources and then many more to verify a certain behavior, making regression testing and verification tests complicated. To address this, we have extended Brunet to support a simulation mode, where time is based upon a virtual timer incremented based upon an event driven scheduler. Furthermore, the simulation framework also uses a specialized transports layer to avoid the overhead of using TCP or UDP on the host system, both of which are limited resources and can hamper the ability to simulate large systems. In addition, they offer few if any additional insights into system performance. The specialized framework uses datagrams to pass messages between nodes, thus from the nodes perspective, it is very similar to a UDP transport and can simulate both latency and packet dropping. In this paper, we set the latency between all node pairs to 100 ms.

Both simulation and real system evaluation provide unique advantages. Simulations allow faster than real time execution of reasonable sized networks (up to a few thousand) using a single resource, while enabling easy debugging. Where as deployment on real systems presents opportunities to add more chaos into the system which can be difficult to duplicate in simulation, such as occasional network glitches, CPU delays on processing, and other random difficult to ascertain events.

For the purpose of this paper, we will perform experiments using both deployed systems either on local resources and PlanetLab or using simulation. The various evaluations dictate, which mechanism we use. Though when possible, we evaluate with as many as possible, thus verifying the system.

IV. TOWARDS PRIVATE OVERLAYS

Many users of IPOP began by trying the shared overlay and, once comfortable, attempt to host their own infrastructure. Some are successful without assistance from us, while a majority are not. The most common issue preventing users

from hosting their own independent IPOP systems was the result of network configuration issues. In short, users were able to easily join the shared overlay, but similar attempts to construct their own were hindered.

The motivation for a private IPOP overlay is quite clear, a shared overlay provides no means for controlling access and can easily be tampered. Though bootstrapping a P2P system requires expertise in network administration. To enable users to bootstrap their own private overlays, we previously investigated means by which a public overlay could be used to bootstrap a private overlay.

Our approach for bootstrapping private systems requires an overlay to support a methods for peers to discover each other, relay messages, and obtain their public address mapping as described in our recent paper [12], that focuses on bootstrapping P2P systems. Examples of other potential bootstrap overlays include popular and well established P2P systems, such as Gnutella, Skype, and Kademia. Our initial work supports bootstrapping from XMPP (Jabber) systems and our own P2P overlay, Brunet.

To bootstrap from an existing Brunet overlay, peers first insert their public overlay address into the key represented by $\text{hash}(\$PrivateOverlayNamespace)$ and continue to do so regularly until they disconnect, so as to not let the entry become stale and disappear. Peers attempting to bootstrap into the private overlay can then query this key and obtain a list of public overlay nodes that are currently acting as proxies into the private overlay. By using the public overlay as a transport, similar to UDP or TCP, the private overlay node forms bootstrapping connections via the public overlay. At which point, overlay bootstrapping proceeds as normal. The entire process is represented in Figure 1.

In a small private overlay, there is a reasonable chance that not a single node in that overlay has a public address, making it difficult for the overlay to provide its own form of NAT

traversal services. Rather than having a special case for NAT traversal for the private overlays that differentiates from the public overlay it bootstrapped from, the two share underlying TCP and UDP sockets. This mechanism, known as pathing, allows a single UDP socket and listening TCP socket to create links for many overlays. This is only possible due to the generic transports library, which does not differentiate UDP, TCP, or even relayed links. Thus during link establishment, the pathing system acts as a proxy, by intercepting a link creation request from a specific entity, mapping that to a path, and then requesting from the remote entity a link for that path. The underlying link is then wrapped by pathing and given to the correct overlay node. Resulting in a completely transparent multiplexing of a TCP and UDP socket enabling the NAT traversal in one overlay to benefit the other. Furthermore, once a link has been established, the pathing information is irrelevant, limiting the overhead into the system to a single round trip time in the bootstrapping phase.

A. Time to Bootstrap a Private Overlay

In our previous work [12], we only verified that private overlays could be bootstrapped using a small set of resources, 5, behind NATs. The focus of this paper extends beyond that paper and emphasizes overheads. This evaluation uses both PlanetLab and simulation to determine the delay in bootstrapping an IPOP-only overlay from a public shared overlay in contrast to a common IPOP overlay. 100 tests were run for each of the various network sizes. In the case of simulation, a new network was generated for each test, while PlanetLab reused the same overlay, but the joining peer changed its P2P address as to cause the peer to connect to a different set of nodes in the public overlay. The results are presented in Figure 2.

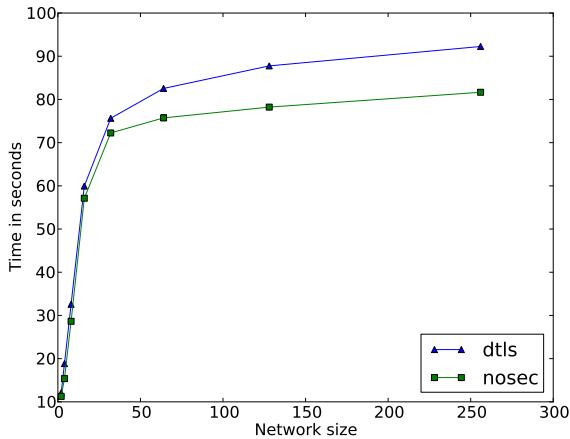


Fig. 2.

There are two important details that we would like to share. First, the PlanetLab side of the experiment is extremely limited in the smaller sized networks, as the system stayed static and a poorly connected peer can shift the results significantly. We have only included the results to show that the techniques do

scale for large systems. Second, simulation of the bootstrapping, we discovered that use the simplistic algorithm described earlier sometimes would result in partitioned overlays. The cause was due to peers determining from a random subset of the peers they initiated connections with that they were connected to the complete overlay, but in reality, they were forming partitions. Because the peers thought they were fully connected, the overlays would never fix the partitioning. To address this, we have set it so that peers occasionally query the DHT entry to determine if there is a peer in the list that should be directly connected to them but is not. If such a peer exists, the searching peer begins connection procedures with that peer.

B. Overhead of Pathing

Much like the previous experiment, this one too is used to verify the use of the private overlay bootstrapping components for VPN usage. In this case, the pathing technique is evaluated to determine what overheads may exist by using the approach. To determine the overheads, two IPOP VNs are deployed on resources on the same gigabit LAN using netperf to evaluate both latency and throughput. We ignore the other specifications of the machines as we compare both the system with pathing versus the system without pathing. The results are presented in Table I. The results indicate that the use of pathing presents negligible overhead for both throughput and latency, making it a very worthy trade to transparently deal with NAT and firewall traversal.

	Latency (ms)	Throughput (Mbit/s)
Standard	0	0
Pathing	0	0

TABLE I
PATHING OVERHEADS

V. SECURITY FOR THE OVERLAY AND THE VPN

Structured overlays are difficult to secure and a private overlay does not imply that it is safe from malicious users as it provides no means to limit access to the system. Malicious users can pollute the DHT, send bogus messages, and even prevent the overlay from functioning, rendering the VPN useless. To address this in means that make sense for VPNs and common users, we have employed a public key infrastructure (PKI) to encrypt and authenticate both communication between peers as well as communication between peers through the overlay, called point-to-point (PtP) and end-to-end (EtE) communication, respectively.

The motivation for using a PKI is that users can either pre-exchange public keys through a trusted medium or place their trust into a third party known as a certificate authority (CA). Unlike other security systems, in particular a key distribution center, which relies on a middleman to establish secure sessions, a CA system enables two peers who have previously obtained a CA signed certificate to establish a trusted relationship directly. So not only can peers form a

relationship directly, they can do so without requiring that the CA be online.

The reasons for securing PtP and EtE are different. Securing PtP communication prevents unauthorized access to the overlay, as peers must authenticate with each other for every link created. Though once authenticated, a peer can perform malicious acts and since the overlay allows for routing over it, the peer can disguise the origination of the malicious acts. By also employing EtE security, the authenticity of messages transferred through an overlay can be easily verified. Though EtE security by itself, will not prevent unauthorized access into the overlay. By employing both PtP and EtE, overlays can be secured from uninvited guests from the outside and can easily identify malicious users on the inside. Implementing both leads to important questions: what mechanisms can be used to implement both and what are the effects of both on an overlay and to a VPN on an overlay.

A. Implementing Overlay Security

There are various types of PtP links, for example, there are TCP and UDP sockets, relays across nodes and overlays, and explored in previous work relays across external services like XMPP. EtE communication is strictly datagram unless the overlay employs a userspace stream based reliability service. Traditional approaches of securing communication such as IPsec or SSL will not apply, nor will approaches that rely on reliable, in order connections such as SSL or TLS. As such, we have implemented an abstraction akin to a security filter as presented in Figure 3, which enables the use of security libraries and protocols that are not reliant on strict interfaces. To this date, we have implemented both a DTLS [13] filter using the OpenSSL implementation of DTLS as well as a protocol that reuses cryptographic libraries provided by .NET that behaves similarly to IPsec.

A security filter has two components, the manager and individual sessions or filters. While the individual sessions could act as filters by themselves, by combining with a manager, they can be configured for a common purpose and security credentials. This approach enables the use of security to be transparent to the other components of the system as the manager handles session establishment, garbage collection of expired sessions, and revocation of peers.

The strength in using certificates is the embedding of identity in the certificate. In network systems, the certificate uses the domain name to uniquely identify and limit the use of a certificate. When a CA signs the certificate, by including the domain name, it ensures that users can almost blindly trust that a certificate is valid, while used to secure traffic to that domain. Communication with another domain using the same certificate will raise a flag and will result in the user not trusting the certificate. In environments with NATs, dynamic IP addresses, or portable devices, typical of P2P systems, assigning a certificate to a domain name will be a hassle as it contrains mobility and the type of users in the system, furthermore, most users are unaware of their IP address and changes to it. Instead, a certificate are signed

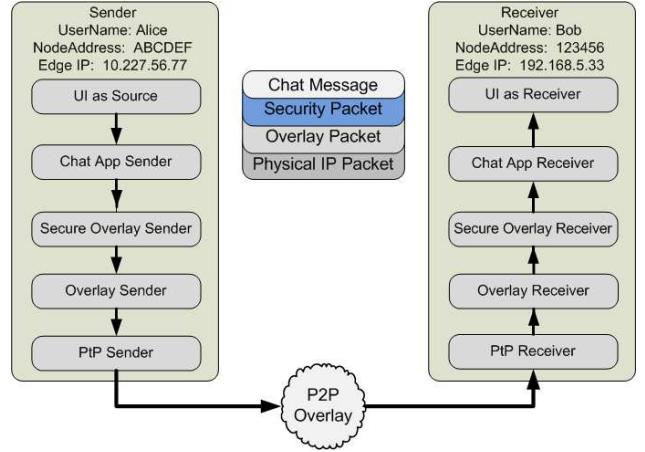


Fig. 3. An example of the abstraction of senders and receivers using a EtE secured chat application. Each receiver and sender use the same abstracted model and thus the chat application requires only high-level changes, such as verifying the certificate used is Alice's and Bob's, to support security.

against the users P2P address. During the formation of PtP links or while parsing EtE messages, the two nodes discover each others P2P addresses, if the addresses do not match the address on the verified certificate, the communication need not proceed further.

The use of a filter approach, this requires one change to the core software, such that it verifies the identity of the remote side prior to allowing packets to traverse the session. In our system, we did this by means of a callback, which presents the underlying sending mechanism, EtE or PtP, and the overlay address stored in the certificate. The receiver of the callback can attempt to cast it into known objects, if successful, it will compare the overlay address with the sender type. If unsuccessful, it ignores the request. If any callbacks return that the sender does not match the identifier, the session is immediately closed.

The last consideration comes in the case of EtE communication that provides an abstraction layer. For example, in the case of VPNs, where where a P2P packet contains an IP packet and thus a P2P address maps to a VPN IP address, a malicious peer may establish a trusted link, but then hijack another users IP session. As such, the application must verify that the IP address in the IP packet matches the P2P address of the sender of the P2P packet.

B. Overheads of Overlay Security

When applying an additional layer to a P2P system, there will be obvious overheads in terms of time to become connected to the overlay. Other less obvious effects will be throughput, latency, and processing overheads, this is primarily due to the use of the P2P system over a wide area network, where the latency and throughput limitations due to the network conditions between two points will make the overhead of security negligible except in the case of bootstrapping due to additional round trip messages used for forming a secure connection.

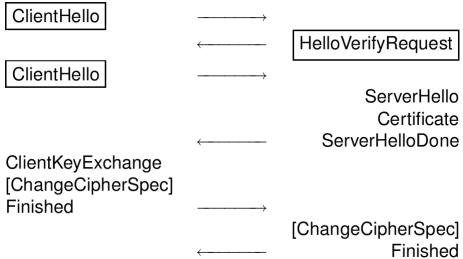


Fig. 4. DTLS handshake

For example, consider the DTLS handshake as presented in Figure 4, which consists of 6 messages or 3 round trips. PtP security may very well have an effect on the duration of overlay bootstrapping. There even exists a possibility that with more messages during bootstrap, the probability one drops is higher, which could have also, in turn, have an effect, though possibly negligible, on time to connect. To evaluate these concerns, we have employed both simulation and real system experiments.

The following experiments use both the simulation to evaluate time to connect a new node to an existing resource. Then another experiment is performed to evaluate how long it takes to bootstrap various sized overlays if all nodes join at the same time. This experiment is only feasible via simulation as attempting to reproduce in a real system is extremely difficult due how quickly the operations complete.

1) Adding a Single Node: This experiment determines how long it takes a single node to join an existing overlay with and without DTLS security. The experiment is performed using both simulation and a real deployed system. After deploying a set of nodes without security and with security on PlanetLab, a crawl was run to determine the size of the network. In both cases, the overlay maintained an average size of around 600 nodes. At which point, a new node connected to the overlay 1,000 times, generating a new P2P address each time, and thus connecting to a different point in the overlay. As soon as the node believed it had connected to the correct location in the overlay, the experiment ended. For simulation purposes, a new overlay was deployed each time, 100 times, each time a new node would join. The test completed after the entire overlay agreed upon being properly correct, a more rigorous test than that tested in the real system. The CDF for the experiments is presented in Figure 5.

2) Bootstrapping an Overlay: The purpose of this experiment is to determine how quickly an overlay using DTLS can bootstrap in comparison to one that does not given that there are no existing participants. Nodes in this evaluation are randomly given information about 5 different nodes in the overlay and then all attempt to connect with each other at the same time. The evaluation completes after the entire overlay has all nodes connected and in their proper position. For each network size, the test is performed 100 times and the average result is presented in Figure 6.

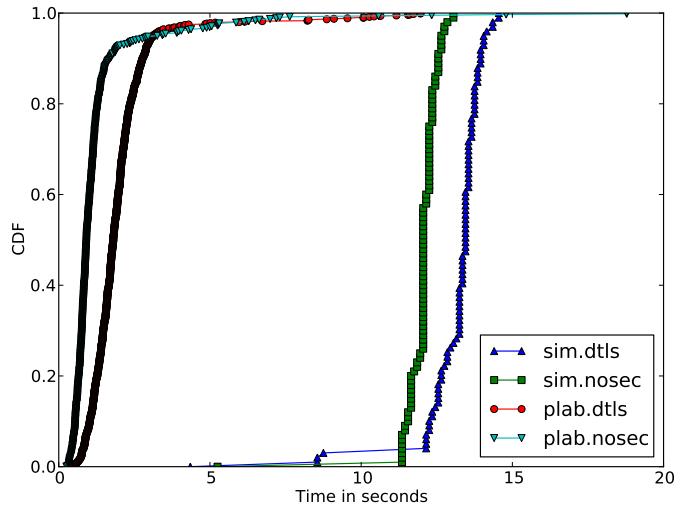


Fig. 5. Time in seconds for a single node to join a secure (dtls) and insecure (nosec) structured overlay, using both PlanetLab (plab) and the Simulator (sim).

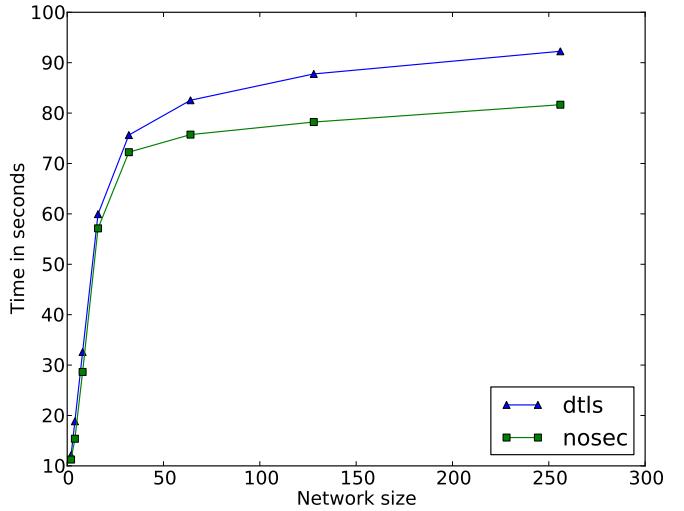


Fig. 6. Time in seconds for a secure (dtls) and insecure (nosec) structured overlay to bootstrap, given that all nodes bootstrap simultaneously.

C. Discussion

Both evaluations show that the overhead in using security is practically negligible, when an overlay is small. In the case of adding a single node, it is clear that the simulation and deployment results agree, as the difference between bootstrapping into an overlay with and without security remains nearly the same. Clearly this motivates the use of security if time to connect is the most pressing question.

While staring at the graphs, we were quite surprised by how little the time to bootstrap a secure overlay was than an insecure overlay. What we realized is that complex connection handshaking, as implemented in Brunet, can easily dominate connection time. For example, in Brunet, two peers must communicate via the overlay prior to forming a connection, and the system differentiates between bootstrapping connections

and overlay connections. Thus even though a peer may have a bootstrapping connection, it will need to go through the entire process to form an overlay connection with a peer. The decision for this process was to keep the code simple and easier to maintain.

VI. HANDLING USER REVOCATION IN A SECURE OVERLAY

Unlike decentralized systems that use shared secrets, in which the creator of the overlay becomes powerless to control malicious users, a PKI enables the creator to effectively remove malicious users. Typical PKIs either use a certificate revocation list (CRL) or online certificate verification protocols such as Online Certificate Status Protocol (OCSP). These approaches are orthogonal to decentralized systems as they require a dedicated service provider in order to verify a certificate, if the service provider is offline, an application can only rely on historical information to make a decision on whether or not to trust a link. In a decentralized system, these features can be enhanced so not to rely on a single provider. In this section, we present two mechanisms of doing so: storing revocations in the DHT and performing overlay broadcast based revocations.

A. DHT Revocation

A DHT can be used to provide revocation similar to that of OCSP or CRLs, where users can either query a service provider to obtain the validity of one or more certificates. A revocation can be stored in the DHT by signing the hash of a certificate and the time stamp of revocation and storing all three in the DHT at the key formed by the hashing of the certificate. In doing so, revocations will be uniformly distributed across the overlay, not relying on any single entity.

The problem with the DHT approach is that it does not provide an event notification for members currently communicating with the peer. While peers could continue to poll the DHT to determine a revocation, doing so is inefficient. Furthermore, a malicious peer, who has a valid but revoked certificate could force every member in the overlay to query the DHT, negatively affecting the DHT nodes storing the revocation.

B. Broadcast Revocation

Broadcast revocation can be used to address the deficiencies of DHT revocation. As a topic of previous research works [14], [15], structured overlays can be used without additional state to perform efficient broadcasts from any point in the overlay to the entire overlay. The form of broadcast can be used to perform to notify the entire overlay immediately about a new revocation. In these papers, analysis and simulations have shown that the approach can be completed in $O(\log^2 n)$ time.

Our modified algorithm as illustrated in Figure 7 utilizes the organization of a structured system with a circular address space that requires peers be connected to those whose node addresses are the closest to their own, features typical of 1-d structured overlays including Chord [16], Pastry [17], and

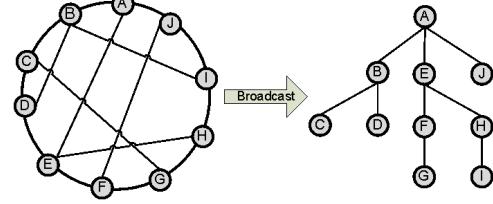


Fig. 7. Broadcast performing a complete overlay broadcast

Symphony. Using such an organization, it is possible to do perform a broadcast with no additional state. To perform a broadcast, each node performs the following recursive algorithm:

```
BROADCAST(start, end, message):
RECEIVE(message)
for  $i$  in length(connections) do
     $n\_start \leftarrow \text{ADDRESS}(\text{connections}[i])$ 
    if  $n\_start \notin [\text{start}, \text{end}]$  then
        continue
    end if
     $n\_end \leftarrow \text{ADDRESS}(\text{connections}[i + 1])$ 
    if  $n\_end \notin [\text{start}, \text{end}]$  then
         $n\_end \leftarrow \text{end}$ 
    end if
     $\text{msg} \leftarrow (\text{BROADCAST}, n\_start, n\_end, \text{message})$ 
     $\text{SEND}(\text{connections}[i], \text{msg})$ 
end for
```

with “connections” as a circular list of connections in non-decreasing order from the perspective of the node performing the current recursive, broadcast step.

In this algorithm, broadcast initiator uses its own address as the start and end, thus the broadcast will span the entire overlay after completing recursive calls at each connected node. A recursive end, “ n_end ”, must be inside the region between “start” and “end”, thus if the connection following the current sending connection, “ $\text{connections}[i + 1]$ ”, is not in that region, it will only broadcast up to “end” and not the address specified by that connection. Finally, nodes, who have a connection to the malicious peer, will end the connection prior to accidentally forwarding the message to the peer by receiving and acting upon the revocation prior to forwarding the message. To summarize, the overlay is recursively partitioned amongst the nodes at each hop in the broadcast. By doing so, all nodes receive the broadcast without receiving duplicate broadcast messages.

C. Evaluation of Revocation Models

D. Discussion

In contrast to the DHT solution, broadcast revocation occurs only once and does not leave state behind. Thus the broadcast is not a complete solution, new peers to the overlay or those who missed the broadcast message will be unaware of a revocation. Furthermore, if an overlay is shared by many VPNs, it may prevent overlay broadcasting or itself may be inefficient.

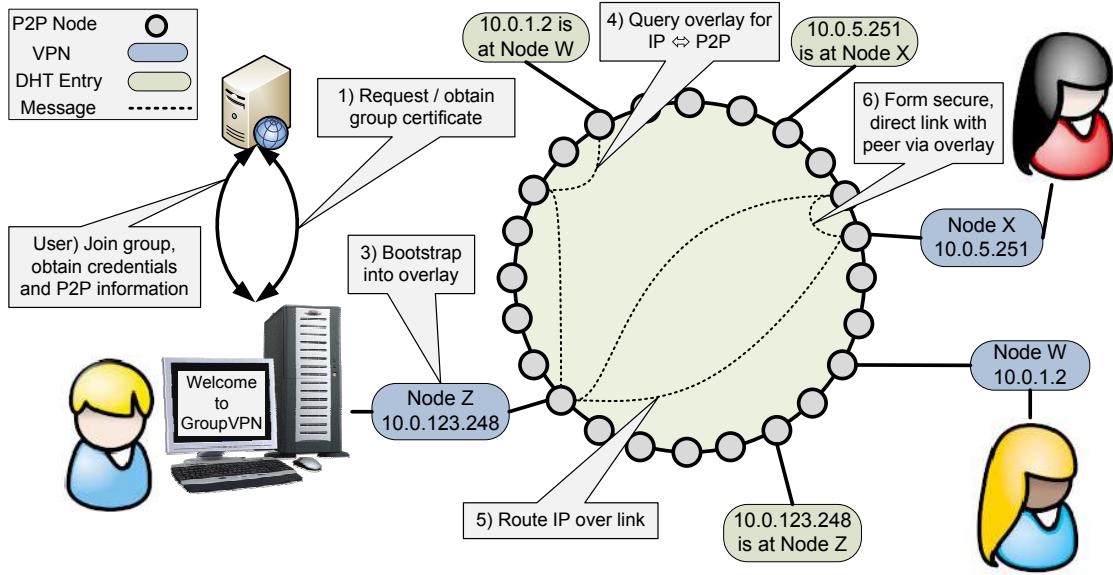


Fig. 8. Process in bootstrapping a new GroupVPN instance.

The DHT solution by itself may also not sufficient as revocations may be lost over time as the entries must have their leases renewed in the DHT. To address this condition, each peer maintains a local CRL and the owner of the overlay can occasionally send updates to the CRL through an out of band medium, such as e-mail. A better long term solution may be the use of a gossip protocol so that peers can share their lists with each other during bootstrapping phases.

A key assumption in using these is that a Sybil [18], or collusion attack, is difficult in the secured overlay. If a Sybil attack is successful, both a DHT and broadcast revocation may be unsuccessful, though peers could fix this problem by obtaining the CRL out of band. In addition, previous work [19] has described decentralized techniques to limit the probability of such attacks from occurring. In our approach, the use of central authority to review certificate requests can be used to limit a single user from obtaining too many certificates as well as ensuring uniform distribution of that user's P2P addresses, further hampering the likelihood of a Sybil attack. The ability to automate this is left as future work.

VII. MANAGING AND CONFIGURING THE VPN

While the PKI model applies cleanly to P2P models, setting up, deploying, and then maintaining security credentials can easily become a non-negligible task, especially for non-experts. Most PKI-enabled systems require the use of command-line utilities and lack methods for assisting in the deployment of certificates and policing users. In order to facilitate use in real systems with non-experts, it is important to have an easy to use framework. Our focal solution to this issue is a partially automated PKI reliant on a group based web interface distributable in forms of Joomla add-ons as well as a virtual machine appliance. In this environment, groups can share a common site with each group having their own unique CA. Although this does not preclude other methods of

CA interaction, experience has shown that it provides a model that is satisfactory for many use cases.

A group based Web 2.0 environment enables low overhead configuration of collaborative environments. The roles in a group environment can be divided into administrators and users. Users have the ability to join and create groups; whereas administrators define network parameters, accept or deny join requests, remove users, and promote other users to administrators. By applying this to a VPN, the group environment provides a wrapper around a public key infrastructure (PKI), where the administrators of the group act as the certificate authority (CA) and the members have the ability to obtain signed certificates. Elaborating further, when a user joins a group, the administrator can enable automatic signing of certificates or require prior review; and when peers have overstayed their welcome, an administrator can revoke their certificate by removing them from the group. Revocations are stored on the site as a CRL (certificate revocation list) and when they occur are broadcast onto the overlay and inserted into the DHT of the respective overlay. Though for these forms of systems a user revocation list as opposed to a CRL simplifies revocation, since users and not individual certificates will be revoked.

Administrators of a group configure the VPN address range, namespace, security, and the ability to specify reuse of an existing overlay or a list of user managed nodes. When a user has been accepted into the group, they are able to download VPN configuration data, which can be seamlessly added to the VPN by running the VPN configuration process. In addition to IP address range, namespace, and security options, the configuration data also contains the group website's address and a shared secret. The shared secret uniquely identifies the user, so that the website can automatically sign the certificate or enqueue it so the administrator can manually sign it. Certificate requests consist of a public key and the user's

shared secret and are sent over HTTPS to the web server. The website creates and signs a certificate request based upon the public key and the user's relevant information ensuring that users cannot trick the website into signing malicious data. Upon receiving the signed certificate, peers are able to join the private overlay and VPN enabling secure communication amongst the VPN peers. The entire bootstrapping process including address resolution and communication with a peer is illustrated in Figure 8.

There are many ways of implementing and hosting the web site. For example, Google offers free hosting of Python web applications through Google Apps, but this requires that the user owns a domain. Alternatively, the user could host the group site on a public VN, peers interacting with the GroupVPN would need to connect with the public VN in order to create an account, get the configuration data, and to sign their certificate at which point they could disconnect from it. This does not preclude the use of other social mediums nor a central site dedicated to the formation of many GroupVPNs. Many GroupVPNs can share a single site, so long as the group members trust the site to host the CA private key.

VIII. RELATED WORK

A. VPNs

Hamachi [20] is a centralized P2P VPN provider. Peers authenticate with the website, use it to discover remote peers, and then connect with them. The website itself provides a central group management environment, which can be used to control access to the VPN. While the Hamachi protocol claims to support various types of security [21], the implementation appears to only support the key distribution center (KDC). This requires that all peers establish trusted relationship through the central website. The Hamachi approach makes it easy for users to deploy their own services, but places limitations on network size, costs money, and does not allow users to fork off from Hamachi and maintain their own system, if they so desire. In contrast, our approach presents a completely decoupled environment, which allows peers to start using our shared system to bootstrap private overlays and migrate away without cost if need be. Our approach also has much higher uptimes as the only process that is centralized is the signing of the certificate, whereas if the Hamachi central server goes offline, no new peers can join the VPN.

Campagnol VPN [22] provides similar features to Hamachi, that is a P2P VPN that relies on a central server for rendezvous or discovery of peers. The key differences between Hamachi and Campagnol is that Campagnol is free and does not provide a service, users must deploy their own rendezvous service. The authors of Campagnol also state that the current approach limits the total number of peers sharing a VPN to 100 so not to overload the rendezvous service. The current implementation does not support a set of rendezvous nodes, though doing so would make the approach much more like ours. In addition, the system relies on traditional distribution of a CRL to handle revocation.

The last VPN we discuss is Tinc [23]. Tinc is a mature, open source VPN, dating back to the early part of 2000. Tinc is a decentralized VPN, which requires that users manually organize an overlay and the Tinc uses its own routing protocols to determine optimum paths. In comparison to our approach, Tinc cannot be used to allow peers to bootstrap private VPNs from an existing VPN nor does Tinc automatically handle churn in the VPN. If node connecting two separate pieces of the VPN overlay goes offline, the VPN will be partitioned until a user manually creates a link connecting the pieces. Furthermore, Tinc does not form direct connections for latency and throughput reasons, thus members acting as routes in the overlay incur the price of acting as packet forwarders.

B. P2P Systems

BitTorrent [24], a P2P data sharing service, supports stream encryption between peers sharing files. The purpose of BitTorrent security is to obfuscate packets to prevent traffic shaping due packet sniffing. Thus BitTorrent security uses a weak stream cipher, RC4, and lacks peer authentication as symmetric keys are exchanged through an unauthenticated Diffie-Hellman process.

Skype [25] provides decentralized audio and video communication to over a million concurrent users. While Skype does not provide documentation detailing the security of its system, researchers [26], [27] have discovered that Skype supports both EtE and PtP security. Though similar to Hamachi, Skype uses a KDC and does not let users setup their own systems.

As of December 2009, the FreePastry group released an SSL enabled FreePastry [17]. Though relatively little is published regarding their security implementation, the use of SSL prevents its application for use in the overlay and for overlay links that do not use TCP, such as relays and UDP. Thus their approach is limited to securing environments that are not behind NATs and firewalls that would prevent direct TCP links from forming between peers.

C. Certificate Authorities

The RobotCA [28] provides an automated approach for decentralized PKI. A RobotCAs receives request via e-mail, verifies that the sender's e-mail address and embedded PGP key match, signs the request, and mails it back to the sender. RobotCAs are only as secure as the underlying e-mail infrastructure and provide no guarantees about the person beyond their ownership of an e-mail address. A RobotCA does not provide features to limit the signing of certificates nor does it provide user-friendly or intuitive mechanisms for certificate revocation.

IX. CONCLUSIONS

This paper describes a novel that verifies the assertion that structured overlays can be used to design elegant VPNs. Explicitly, we have taken common structured overlays and explored organization, public overlays for connectivity, and private overlays for security and then described our GroupVPN

which binds them the components together to create collaborative environments for configuration and management of VPNs. This paper extends our previous work, IPOP, a P2P virtual network system, to support user-friendly approaches for users to create and manage their own IPOP systems with security. To do this, each IPOP system bootstraps into its own unique, secure P2P overlay. This approach not only enables significantly more secure IPOP deployments but also enables for more efficient overlay multicast and broadcast and the cost of doing so amounted to only a few hundred lines of code.

The use of service overlays significantly improves performance and maintenance. Peers can easily control membership in the overlay and it presents unique opportunities for decentralized revocation. A DHT approach allows results to be stored on the overlay instead of using centralized CRLs and broadcast to immediately notify active participants of a revocation. For future work, we plan further investigating security concerns of using the decentralized revocation techniques and hope to do away with the need of any form of dedicated CRL. Furthermore, we plan on investigating the use overlay broadcasting for IP broadcasting and multicasting, though the current approach places an unfair burden on the first few hops of a broadcast.

Without the functionality of GroupVPN projects like Archer [1], would be impractical. Archer consists of over 500 resources from 5 different universities, including University of Florida, Florida State University, Northeastern University, University of Minnesota, and University of Texas. In the past year, since Archer came online, over 100 unique users have contributed and taken advantage of the voluntary computing cycles. A new user to the system begins by creating an account at Grid-Appliance.org and requesting membership in the Archer GroupVPN group. Once access has been granted, users can obtain configuration data used by the Grid Appliance initialization scripts to seamlessly add resources to the grid. This method allows independent submission sites, unlike most grid systems that have a shared submission site, which require dedicated administrators. Most users connect to the system using a pre-configured virtual machine appliance, so that they do not need to be experts in grid systems to take advantage of Archer. Enabling this using decentralized VPNs would be difficult as the user would need to create manual links to the rest of the system for each new resource. N2N may work, but the network size of Archer is larger than the recommendations made by N2N and would still require the setup of address allocation facilities. In general, all existing approaches would fail besides those with centralized components, because, at the time of this writing, all of Archer's resources are behind NATs. Even though centralized could be used, it would require additional dedicated resources and management, limiting access if the central component went offline.

The GroupVPN has been used as the virtual network for the Grid Appliance, which is the basis of Archer and, in general, enables the creation of decentralized, collaborative environments for computing grids. Recently, grids at La Jolla Institute for Allergy and Immunology and two in Eastern

Europe went live using GroupVPN without receiving any technical support from us. Researchers at Clemson University and Purdue have opted for this approach over centralized VPNs as the basis of their future distributed compute clusters and have actively tested networks of over 700 nodes.

REFERENCES

- [1] R. J. Figueiredo, P. O. Boykin, J. A. B. Fortes, T. Li, J. Peir, D. Wolinsky, L. K. John, D. R. Kaeli, D. J. Lilja, S. A. McKee, G. Memik, A. Roy, and G. S. Tyson, "Archer: A community distributed computing infrastructure for computer architecture research and education," in *CollaborateCom*, November 2008.
- [2] A. Ganguly, A. Agrawal, O. P. Boykin, and R. Figueiredo, "IP over P2P: Enabling self-configuring virtual IP networks for grid computing," in *International Parallel and Distributed Processing Symposium*, 2006.
- [3] M. Castro, P. Druschel, A.-M. Kermarrec, and A. Rowstron, "One ring to rule them all: Service discover and binding in structured peer-to-peer overlay networks," in *SIGOPS European Workshop*, Sep. 2002.
- [4] M. Castro, M. Costa, and A. Rowstron, "Debunking some myths about structured and unstructured overlays," in *NSDI'05: Proceedings of Symposium on Networked Systems Design & Implementation*.
- [5] P. O. Boykin, J. S. A. Bridgewater, J. S. Kong, K. M. Lozev, B. A. Rezaei, and V. P. Roychowdhury, "A symphony conducted by brunet," <http://arxiv.org/abs/0709.4048>, 2007.
- [6] G. S. Manku, M. Bawa, and P. Raghavan, "Symphony: distributed hashing in a small world," in *USITS*, 2003.
- [7] J. Rosenberg, "Interactive connectivity establishment (ICE): A protocol for network address translator (NAT) traversal for offer/answer protocols," <http://tools.ietf.org/html/draft-ietf-mimusic-ice-19>, October 2008.
- [8] A. Ganguly, D. Wolinsky, P. Boykin, and R. Figueiredo, "Decentralized dynamic host configuration in wide-area overlays of virtual workstations," in *International Parallel and Distributed Processing Symposium*, March 2007.
- [9] I. Stoica, D. Adkins, S. Zhuang, S. Shenker, and S. Surana, "Internet in-direction infrastructure," *IEEE/ACM Transactions on Networking*, 2004.
- [10] D. I. Wolinsky, Y. Liu, P. S. Juste, G. Venkatasubramanian, and R. Figueiredo, "On the design of scalable, self-configuring virtual networks," in *IEEE/ACM Supercomputing 2009*, November 2009.
- [11] B. Chun, D. Culler, T. Roscoe, A. Bavier, L. Peterson, M. Wavrzonik, and M. Bowman, "Planetlab: an overlay testbed for broad-coverage services," *SIGCOMM Comput. Commun. Rev.*, 2003.
- [12] D. I. Wolinsky, P. St. Juste, P. O. Boykin, and R. Figueiredo, "Addressing the P2P bootstrap problem for small overlay networks," in *10th IEEE International Conference on Peer-to-Peer Computing*, 2010.
- [13] E. Rescorla and N. Modadugu. (2006, April) RFC 4347 datagram transport layer security.
- [14]
- [15]
- [16] I. Stoica, R. Morris, D. Liben-Nowell, D. R. Karger, M. F. Kaashoek, F. Dabek, and H. Balakrishnan, "Chord: a scalable peer-to-peer lookup protocol for internet applications," vol. 11, no. 1, 2003.
- [17] A. Rowstron and P. Druschel, "Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems," in *International Conference on Distributed Systems Platforms*, November 2001.
- [18] J. R. Douceur, "The sybil attack," in *IPTPS '01: Revised Papers from the First International Workshop on Peer-to-Peer Systems*. Springer-Verlag, 2002, pp. 251–260.
- [19] M. Castro, P. Druschel, A. Ganesh, A. Rowstron, and D. S. Wallach, "Security for structured peer-to-peer overlay networks," in *Symposium on Operating Systems Design and Implementaion (OSDI'02)*, December 2002.
- [20] LogMeIn. (2009) Hamachi. <https://secure.logmein.com/products/hamachi2/>.
- [21] LogMeIn, Inc. (2009) LogMeIn hamachi2 security.
- [22]
- [23] G. Sliopen. (2009, September) tinc. <http://www.tinc-vpn.org/>.
- [24] (2007, December) Message stream encryption. http://www.azureuswiki.com/index.php/Message_Stream_Encryption.
- [25] S. Limited. Skype. <http://www.skype.com>.
- [26] D. Fabrice. (2005, November) Skype uncovered. http://www.ossir.org/windows/supports/2005/2005-11-07/EADS-CCR_Fabrice_Skype.pdf.

- [27] S. Guha, N. Daswani, and R. Jain, "An experimental study of the skype peer-to-peer voip system," in *IPTPS'06*, 2006.
- [28] (2005, October) RobotCA. <http://www.wlug.org.nz/RobotCA>.