

On the Design and Implementation of Autonomic, Decentralized GroupVPNs

David Isaac Wolinsky, P. Oscar Boykin, Renato Figueiredo
University of Florida

Abstract—Virtual private networks (VPNs) enable safe and simple use of network applications in distributed, insecure, and constrained environments by providing all-to-all connectivity in a secure, isolated virtual environment. Existing systems present challenges to users lacking networking, security, and operating expertise in small and medium groups, such as academic, business, and home environments. In this paper, we describe a VPN architecture motivated by the challenges with existing systems and their application in these environments. Our contribution is a system that provides autonomic, decentralized VPNs using common Web 2.0 group collaboration systems for configuration and management. To evaluate these claims, we provide a reference implementation and compare it to existing VPN approaches.

I. INTRODUCTION

A Virtual Private Network (VPN) provides the illusion of a local area network (LAN) spanning a wide area network (WAN) infrastructure by creating encrypted and authenticated, secure¹ communication links amongst participants. Common uses of VPNs include secure access to enterprise network resources from remote/insecure locations, connecting distributed resources from multiple sites, and establishing virtual LANs for multiplayer video games and media sharing over the Internet.

The architecture described in this paper addresses usage scenarios where VPNs are desired but complexity in deployment and management limits their applicability, such as collaborative academic environments linking individuals spanning multiple institutions, where coordinated configuration of network infrastructure across the different sites is often impractical. Another example is the small/medium business (SMB) environment, where it is often desirable to interconnect desktops and servers across distributed sites and secure traffic to enterprise networked resources without incurring the complexity or management costs of traditional VPNs. Alternatively, use of a VPN across an extended family enables sharing of media, such as family videos and pictures, though none of the individual sites have the capability of hosting a centralized service nor want computers left on all the time.

In this paper, we discuss an architecture that deals with the following challenges:

- 1) **Configuration:** How will users connect to the VPN, what type of security credentials will be used, and what will the network parameters be?

¹For the remainder of this paper, unless explicitly stated otherwise, security implies encryption and mutual authentication between peers.

- 2) **Peer Management:** How are new users added to the system, what if they have many resources to add as well? What happens when a user becomes malicious, how are users revoked?
- 3) **Connectivity:** Given a hostname for a remote node, how will the VPN discover the remote peer and what will the method of communication be?
- 4) **Privacy:** Will communication in the VPN be transmitted in clear text amongst members but secure in transit or will secure links be established between source and destination?

Centralized approaches (e.g. OpenVPN [1]) require dedicated infrastructures, which by default see all messages in the clear text, and prevent direct communication between peers, limiting bandwidth and creating inefficient multicast paths. Centralized systems with P2P (peer-to-peer) links (e.g. Hamachi [2], Wippien [3], Gbridge [4], PVC [5], allow secure and direct communication amongst peers but rely on a central resource to provide session establishment and are thus vulnerable to man-in-the-middle attacks. Decentralized approaches (e.g., ViNe [6], Violin [7], VNET [8], tinc [9]) require manual configuration of links between members of the virtual network and do not automatically reconfigure when links fail; security in these systems does not secure communication between end hosts. Existing P2P approaches lack scalability and availability (N2N [10]).

In this paper, we present a novel VPN architecture that uses an existing public overlay to self-organize a VPN using a set of P2P systems known as structured overlays and configured with a Web 2.0 group collaborative environment. Popular examples of large, public overlay networks include Gnutella, Kademia, and Skype. P2P systems provide automatic link creation, facilities for discovering peers in an overlay, and efficient methods for broadcasting. We describe methods by which the VPN can be configured from a web interface and use features of P2P systems to create an autonomic and decentralized GroupVPN.

Section II begins the paper by discussing the various VPN approaches. The core features of a VPN using a structured overlay are described in Section III. The following sections, IV and V have a detailed focus on P2P network constraints and efficient IP broadcast and multicast. Section VI presents a validation of our approach in comparison to existing approaches through quantitative analysis. We conclude in Section VII by exploring a usage scenario and how the constraints apply to our approach with existing approaches used in a real

environment.

II. VIRTUAL PRIVATE NETWORKS

VPNs can be divided into two components, a local or end point configuration and a network configuration. Most VPNs share a common local configuration and differ significantly in the area of network configuration. Local configuration deals with user or system interaction with the VPN, whereas network configuration tackles with the challenges of infrastructure and peer discovery, connectivity, and security. This section begins by reviewing the fundamental local configuration features and then reviews common network configurations as presented in Table I.

TABLE I
VPN CLASSIFICATIONS

Type	Description
Centralized	Clients communicate through one or more servers which are statically configured
Centralized Servers / P2P Clients	Servers provide authentication, session management, and optionally relay traffic; peers may communicate directly with each other via P2P links if NAT traversal succeeds
Decentralized Servers and Clients	No distinction between clients and servers; each member in the system authenticates directly with each other; links between members must be explicitly defined
Unstructured P2P	No distinction between clients and servers; members either know the entire network or use broadcast to discover routes between each other
Structured P2P	No distinction between clients and servers; members are usually within $O(\log N)$ hops of each other via a greedy routing algorithm; use distributed data store for discovery

A. The Basic Client VPN Configuration

Figure 1 presents an abstraction of the common features found in all VPN clients: a service that communicates with the VPN system, and a virtual network (VN) device for host integration. During initialization, the VPN service authenticates with the overlay², optionally, querying for information about the network, such as network address space, address allocations, and domain name service (DNS) servers. At which point, the VPN enables secure communication amongst participants.

A VN device allows applications to communicate transparently over the VPN. By providing mechanisms for injecting incoming packets into and retrieving outgoing packets from the networking stack, a VN device enables the use of common network APIs such as Berkeley Sockets, thereby allowing unmodified applications to work over the VPN. To do this VN devices allow the creation of a virtual network interface providing either a virtual Ethernet or IP device. The most widely available, free VN device is TAP [11]. Most operating systems support TAP either through existing features in the OS

²An overlay in this context refers to the topology formed by VPN nodes including central servers, other VPN clients, or relays.

or through third party drivers, these systems include Windows, Linux, Mac OS/X, BSD, and Solaris.

VN devices can be configured manually through command-line tools or OS' APIs or dynamically by the universally supported dynamic host configuration process (DHCP) [12]. When a VN device obtains an IP address, it triggers the OS to add a new rule to the routing table directing all packets sent to the VPN subnet to the VN device. Packets read from the the TAP device are encrypted and sent to the overlay via the VPN client. The overlay delivers the packet to another client or a server with a VN stack enabled. Received packets are decrypted, verified for authenticity, and then written to the VN device. In most cases, the IP layer header remains unchanged, while VPN configuration determines how the Ethernet header is handled.




Fig. 1. A typical VPN client. A VN device makes application interaction with the VPN transparent by integrating with the OS network stack.

B. Centralized VPN Systems

OpenVPN is an open and well-documented platform for deploying centralized VPNs. In this paper, it is used as the basis for understanding centralized VPNs as it represents features common to most centralized VPNs.

In centralized VPN systems, clients forward all VPN related packets to the server. Client responsibilities are limited to configuring the VN device and authenticating with the VPN server; whereas the servers are responsible for authentication and routing between clients. Likewise, broadcast and multicast packets also must pass through the central server.

Centralized systems do not prohibit multiple servers, though each server must be manually configured to interact with remote servers. Clients randomly select a server from a list of known servers, implementing a simple load balance. Once connected, the servers provide the client an IP address in the VPN address space. Thus two peers on different servers will have three hops between each other, two from routing over the two servers and one for the forwarding to the destination.

All inter-client communication flows through a central server. By default, a client encrypts a packet and sends it to the server. Upon receiving the packet, the server decrypts it, determines where to relay it, encrypts it, and then sends the packet to its destination. This model allows a server to eavesdrop on communication. While a second layer of

encryption is possible through a shared secret, it requires out-of-band communication and increases the computing overhead on communication since the packet will now need to be encrypted and authenticated twice at both the VPN source and destination sites.

C. Centralized P2P VPN Systems

Hamachi [2] is the first well-known centralized VPN that used the ambiguous moniker “P2P VPN”. In reality, these systems are better classified as centralized VPN servers with P2P links. Similar VPNs include Wippien [3], Gbridge [4], PVC [5], and P2PVPN³ [13]. The P2P in these systems is limited to direct connectivity between clients orchestrated through a central server: in Wippien it is a chat server, while P2PVPN uses a BitTorrent tracker. If NAT traversal or firewalls prevent direct connectivity, the central server can act as a relay. Each approach uses their own security protocols with most using a server to verify the authenticity and setup secure connections between clients. In regards to the P2PVPN, long term goals involve the creation of an unstructured, which would provide a method of decentralized organization.

D. Decentralized VPN Systems

Some examples of systems that assist in distributing load in VPN systems are tinc [9], CloudVPN [14], ViNe [6], VNET [8], and Violin [7]. These systems are not autonomic and require explicit specification of links between resources. This means that, like OpenVPN, these systems can suffer VPN outages when nodes go offline, thus administrators must maintain the VPN connection table. Unlike OpenVPN, these approaches typically do not require all-to-all direct connectivity for all-to-all communication. Users can either setup out-of-band NAT traversal or route through relays. Links are manually configured.

E. Unstructured P2P VPN Systems

Unlike centralized and decentralized systems, P2P environments require the user to connect to the overlay, which then automatically configures links. The simplest form of overlays are unstructured, where peers form random connections with each other and use broadcast and stochastic techniques to find information and other peers, though due to its unstructured nature, the system cannot guarantee distance and routability between peers. The only example of an unstructured VPN is N2N [10]. In N2N, peers first connect to a super node and then, to find another peer, they broadcast discovery messages to the entire overlay. In the case that peers cannot form direct connection, peers can route to each other over the N2N overlay. In the realm of VPNs, all client VPNs are also servers performing authentication though neither approach deals with decentralized address allocation. N2N does not address the challenges of IP address allocation and management, relegating that to the user to handle.

³Due to the similarities between the name P2PVPN and focus of this paper, “P2PVPN” refers to [13] and “P2P VPN” to the use of P2P in VPNs.

III. STRUCTURED P2P VPN SYSTEMS

This section describes the basic construction of a Structured P2P VPN, including background on structured overlays, address allocation and discovery, security, and connectivity.

A. Background

An alternative to dealing with scalability concerns in unstructured systems, structured overlays provide distributed look up services with guaranteed search time with a lower bound of $O(\log N)$, in contrast to unstructured systems, which rely on global knowledge/broadcasts, or stochastic techniques such as random walks [15]. Some examples of structured systems are Pastry [16], Chord [17], Symphony [18], Kademlia [19], CAN [20], and Brunet [21]. In general, structured systems are able to make these guarantees by self-organizing into a well-defined structured topology, such as a 2D ring or a hypercube, by randomly generated, though uniformly distributed, node identifiers. A key component of most structured overlays is support for decentralized storage/retrieval of information by mapping keys to specific node IDs in an overlay called a distributed hash table (DHT). At a minimum, the data is stored at the node ID either smaller or larger to the data’s node ID and for fault tolerance the data can be stored at other nodes. DHTs can be used by peers of systems to coordinate allocation and discovery of resources, making them attractive for self-configuration in decentralized collaborative environments.

B. Address Allocation and Discovery

In the context of VPNs, structured overlays can handle address allocation and discovery through the use of a DHT, such systems have been proposed in [22], [23]. In order to connect with the VPN overlay, each VPN member becomes a member of the structured overlay. To address the challenges of having multiple VPNs in the same overlay, each VPN has its own namespace, reducing the likelihood of overlap. To enable scalable and decentralized address allocation and discovery, peers store mappings of IP address to node ID into the DHT, typically of the form $\text{hash}(\text{namespace} + \text{IP}) = \text{nodeID}$. Thus a peer attempting to allocate an address will insert this key, value pair into the overlay. The first peer to do this will be the owner of the IP address allocation. Therefore the DHT must either support atomic writes into the DHT and prevent future writes or provide creation times of key, value pairs during queries.

Mechanisms to self-configure the IP address and network parameters of the local system can be provided by DHCP, the user manually configuring the IP address, or the VPN hooking into OS APIs. Address discovery is initiated when an outgoing packet for a remote peer arrives at the VPN software. At which point, the VPN will query the DHT for the node ID who owns the IP address and forward the packet to the destination. More discussion on both these topics is further covered in our previous work [24].

C. Securing the VPN

Structured overlays are difficult to secure. Malicious users can pollute the DHT, send bogus messages, and even prevent the overlay from functioning. When using a structured overlay for a VPN, tamper-resistance and reliable routing mechanisms are ideal, but a structured overlay requires both the point-to-point (PtP), or communication between connected peers, and end-to-end (EtE), or communication between peers across the overlay, security. While tamper resistance and reliability will reduce the ability to affect an overlay, an expectation of a VPN is that no members are malicious, and if a member is found to be malicious their access revoked. Then the question becomes, why so much security? PtP security ensures that malicious third parties cannot insert material into the system, while EtE ensures message security between the communicating peers. Enabling just PtP security would limit the security to the type of VPN security found in other VPNs, all member must be trusted, whereas EtE enables message security even from overlay members.

In our previous work [25], we described mechanism that enable the construction of a private overlay providing both EtE and PtP security. In this system, EtE and PtP communication has been abstracted into a common sender layer. The security component has been instantiated as a filter, enabling it to be placed anywhere in the sending stack and thus securing either EtE or PtP communication without additional complexity. Existing security filter support uses DTLS [26].

D. Connecting to the VPN

Creating and connecting to a private overlay can be difficult when there are no dedicated resources with public addresses, even if the system has many active members behind NATs and firewalls. This is because peers behind NATs and firewalls may be unable to handle incoming connections without the assistance of a third-party. The same work [25] describes how arbitrary public overlays can act as a middle ground to bootstrapping private overlays, using mechanisms unique to these environments to bootstrap into the private overlay. Of course, this mechanism does not preclude the creation of a private overlay without the use of a public bootstrapping mechanism.

In practical implementation, peers first join a public structured overlay using the public overlay's DHT to find other peers. The private overlay node can then either attempt to connect directly with other private overlay nodes using TCP or UDP or use the public overlay as another transport mechanism. We have termed this approach to creating private overlays, virtual private overlays. In all cases, the communication is secured via the EtE and PtP security filters. This approach enables the construction of an overlay consisting entirely of peers behind NATs and firewalls, who would otherwise be unable to form their own overlay.

E. Managing the VPN

A Group based Web 2.0 environment enables low overhead configuration of collaborative environments. The roles in a

group environment can be divided into administrators and users. Users have the ability to join and create groups; whereas administrators can accept or deny join requests, remove users, and promote other users to administrators. To apply this to a VPN, the group environment provides a wrapper around a public key infrastructure (PKI), where the administrators of the group act as the certificate authority (CA) and the members have the ability to obtain signed certificates. Elaborating further, when a user joins a group, the administrator can enable automatic signing of certificates or require prior review; and when peers have overstayed their welcome, they can have their certificates revoked by removing them from the group. As described in [25], revocation can be stored as a CRL on the group site or distributed through broadcast and DHT on the overlay. Though for these forms of systems a user revocation list (URL) as opposed to a CRL simplifies revocation in the case where a malicious user has automatically signed large amounts of certificates.

In the case of the GroupVPN, administrators of a group configure the VPN address range, namespace, and security. Options are also available to specify the reuse of an existing public overlay or a list of user managed nodes. When a user has been accepted into the group, they are able to download VPN configuration data, which can be seamlessly added to the VPN by running the VPN configuration process. In addition to IP address range, namespace, and security options, the configuration data also contains the group website's address and a shared secret. The shared secret uniquely identifies the user, so that the website can determine how to handle certificate requests. When making a certificate request, the user sends over HTTPS a public key and their shared secret, the website creates and signs a certificate request based upon the public key and the user's relevant information ensuring that users cannot trick the website into signing malicious data. Upon receiving the signed certificate, peers are able to join the private overlay and VPN enabling secure communication amongst the VPN peers.

There are many ways of implementing and hosting the web site. For example, Google offers free hosting of Python web applications through Google Apps, but this requires that the user owns a domain. Alternatively, the user could host the group site on a public VN, peers interacting with the GroupVPN would need to connect with the public VN in order to create an account, get the configuration data, and to sign their certificate at which point they could disconnect from it. This does not preclude the use of other social mediums nor a central site dedicated to the formation of many GroupVPNs. As many GroupVPNs can share a single site, so long as the group members trust the site to host the CA private key.

IV. DEALING WITH NATS AND FIREWALLS

As of 2010, the majority of the Internet is connected via Internet Protocol (IP) version 4. A limitation in this protocol is that there are only 2^{32} addresses (approximately 4 billion) available. With the Earth's population of over 8 billion and each individual potentially having multiple devices that have

Internet connectivity, the IPv4 limitation is becoming more and more apparent. To address this issue there have been two approaches: 1) the use of network address translation (NAT) to enable many machines and devices to share a single IP address but preventing bidirectional connection initiation, and 2) IPv6 which supports 2^{128} addresses. Though even when IPv6 becomes the primary Internet protocol, it does not necessarily imply the demise of NAT or firewalls that only allow outbound connections.

Approaches for handling NAT and firewall traversal come from two main branches: hole punching and relaying. When performing hole punching [27], peers simultaneously attempt to create "holes" in their NAT / firewall devices allowing direct TCP or UDP communication with a peer behind another NAT / firewall. This works by confusing the NAT into believing that the peer behind it initiated the connection. TCP NAT hole punching, though, does not work on systems that use stateful firewalls. Relaying, on the other hand, always works as peers behind NATs communicate with a third peer with whom they have direct connectivity, though with the side-effect of additional latency and the cost for the relay of supporting the bidirectional communication.

A. Relays

Centralized and decentralized VPNs do not suffer from connectivity limitations due to NATs and firewalls as all traffic passes through the central server or managed links. In the centralized P2P approaches that lack dedicated servers, there exists no relay implementation. Though in unstructured networks peers can communicate through the overlay when traversal fails. Similarly a structured P2P VPN can use the overlay, though as an overlay grows in size, so does the hop count between peers increasing latency and reducing bandwidth.




Fig. 2. Creating relays when direct connectivity is not possible. Two members, 0000 and ABCD, are unable to form direct connectivity, thus they exchange neighbor information through the overlay and connect to one of each other's neighbors, creating an overlap. The overlap then becomes a relay path (represented by dashed lines) for two-hop communication.

P2P systems can address the connectivity issue through a distributed, autonomic relaying system, whereby peers form overlapping connections with each other's peers, creating two-hop links. When two nodes discover each other via the overlay and attempt to become connected but fail during NAT and firewall traversal, they exchange peer lists through the overlay. Upon receiving this list, the two peers identify the overlap in the neighbor sets to form a two-hop connection. When peers are not located near each other in the overlay, most likely they will have no overlap. This can be addressed by having peers connect to each other's neighbor set proactively creating overlap, as represented in Figure 2.

With the set of neighbors, the peers can include additional data such as latency and connection lifetime to each of the neighbors. When creating overlap or overlap changes, peers review these metrics to decide which subset of the overlap to use as a proxy leaving the remaining peers as reserves.

B. Usefulness of Relays

To determine the value in two-hop overlays as opposed to using the overlay in restrictive NAT and firewall scenarios, this experiment uses an event-driven modeler that reuses the code base of the prototype structured P2P VPN to faithfully implement routing in an overlay as well as latency between peers. Pair-wise latency in this experiment is set using the MIT King Data Set [28], which consists of all-to-all latencies among 1,740 distributed Internet hosts. The modeler measured average overlay, 2-hop, and 1-hop latency between peers. Overlay latency represents the time to route a packet via the overlay. 2-hop latency is based upon the low-latency formed 2-hop connections, where peers route over the overlap with the lowest latency. 1-hop latency is based upon peers forming a direct connection with each other. Only peers that are not directly connected (i.e., have two hops or more) are analyzed, as a single hop would only benefit under triangular inequalities, which are not a consideration in this work.




Fig. 3. A comparison of the average all-to-all overlay routing, two-hop relay, and direct connection latency in a ring based structured overlay.

Results are presented in Figure 3. The network sizes considered begin at 25 nodes because small networks of sizes 20 and under tend to be fully connected in the underlying structured overlay. It is not until the network size expands past 100 and

towards 200 nodes that relays become significantly beneficial. At 100 nodes, there is approximately a 54% performance increase, whereas at 200 there is an 87% increase and it appears to grow proportionately to the size of the pool, demonstrating the performance advantage of two-hop autonomous relays.

In addition, the autonomous relays functionality were verified in a reference implementation of a structured P2P VPN using a real system. The environment consists of PlanetLab as the public overlay and the Archer environment as the VPN environment. PlanetLab [29] provides a set of over 500 distributed computing nodes all with public IP addresses. Archer provides grid computing to computer architecture researchers and consists of over 500 nodes located at various academic sites. To ensure that peers under study formed two-hop connections, a firewall was instantiated preventing direct communication between the two. In tests, it was revealed that the overlay was always able to self-configure a relay; the bandwidth and latency averages and deviations were 2245 Kbit/s \pm 1080 and 58.1 ms \pm 35.5, respectively.

V. EFFICIENT BROADCAST AND MULTICAST

In the original set of structured P2P virtual networks, broadcast and multicast used DHT to keep a list of all the active peers in the VPN. When a peer wanted to send out a broadcast or multicast packet, it would query this list and then send out unicast packets to each peer either directly or via the overlay. Furthermore, as more peers join the system, so does the count of key, value pairs in the DHT. If the peer wishes to send out many multicast or broadcast messages after another, it would either need to query the DHT every time or maintain a cache. In the case of a cache, new peers joining the overlay would need to notify the entire VPN when they joined so that they could be added to caches. The simplistic approach of querying each time does not scale well, while the other approach involving caches can be quite complicated to implement.

With the use of virtual private overlays, broadcast and multicast can be implemented stateless as they can be routed using overlay broadcast. An example of a scalable ring-based structured P2P overlay broadcast algorithm was presented in [25] to broadcast certificate revocation within a private overlay. The algorithm is called bounded-broadcast, because it is capable of broadcasting to a subset of the ring or the entire overlay. Bounded broadcast uses the following recursive algorithm: Begin with node x triggering a broadcast message over the region $[x, y]$. x has F connections to nodes in the range $[x, y]$. Denoting the i^{th} such neighbor as b_i , the node x sends a bounded broadcast over a sub-range, $[b_i, b_{i+1})$, to b_i , except the final neighbor. Differently stated, b_i is in charge of bounded-broadcasting in the sub-range $[b_i, b_{i+1})$. If there is no connection to a node in the sub-range, the recursion has ended. The final neighbor, (b_F) , is responsible for continuing the bounded broadcast from $[b_F, y]$. When a node receives a message to a range that contains its own address the message is delivered to that node and then routed to others in that range. Figure 4 shows how this bounded broadcast forms a

local tree recursively. The time and bandwidth for a bounded broadcast is $O(\log^2 N)$ and $N * \text{packetsize}$ as shown in [30]. To perform a broadcast on the entire overlay, a peer performs the bounded-broadcast starting from its node ID with the end address being the node ID immediately preceding its own in the address space. Though in VPN situations, many peers may already have connections to most if not all of their VPN peers, thus the broadcast algorithm has been modified to route to only connections created for structured overlay purposes and not explicitly only for VPN purposes. Otherwise in many cases, this algorithm will degenerate into one similar to the unicast approach.




Fig. 4. Bounded Broadcast in range $[x, y]$

The validation of broadcasting was done using the same modeler used in the relay evaluation. This evaluation compares broadcasting to a simple unicast mechanism that assumes the peer has already queried the DHT. Both models are evaluated in a private overlay and in a public overlay consisting of 100,000 peers. The VPN size varies from 10 members up to 10,000. It is important to note that the network modeler does not consider host overheads including bandwidth limitations nor simultaneous packet sending, thus the results are from an ideal system. The modeler measures network bandwidth, sending peer bandwidth, and the time to complete the broadcast. The results are presented in Figure ??.




Fig. 5. Bounded Broadcast in range $[x, y]$

The results in the virtual private overlays indicate that in the ideal system, bounded-broadcast can distribute the message nearly as fast as the unicast method though using significantly less bandwidth. Use in the public overlay is not

recommended unless the broadcast message has a significant amount destinations, which typically will not be the case in GroupVPN. In almost all cases, the sending peers has less bandwidth utilization when using broadcast as opposed to unicast. So beyond being stateless, the approach is not only more efficient in terms of bandwidth but all peers receive the message in a reasonable amount of time in comparison to a unicast approach.

VI. EVALUATION OF VPN MODELS

This experiments explores bandwidth and latency in a distributed VPN system to motivate the usage of P2P links in a VPN. The VPNs used are include our prototype (IPOP), OpenVPN, and Hamachi. OpenVPN represents a typical centralized VPN, while Hamachi represents a well-tuned P2P-link VPN. The evaluation was performed on Amazon EC2 using small instance sized Ubuntu i386 instances to create various sized networks ranging from 1 to 32. OpenVPN uses an additional node as the central server and Hamachi has an upper bound of 16 due to limitations in the Linux version at the time of this evaluation. To perform bandwidth tests, the instances are booted and query an NFS for the list virtual IP addresses, peers are ordered such that half the peers are act as clients and the other half the peers creating a 1 to 1 mapping between all sets. Latency and bandwidth tests are performed using netperf's request-reply and streaming tests respectively. Prior to the start of the tests, peers have no knowledge of each other, except the virtual IP addresses, thus connection startup costs are included in the test. Test are run for 10 minutes diluting the connection initiation overhead but represent an example of real usage. Results from the clients are polled at all locations and averaged together, though the OpenVPN server is measured separately. IPOP and OpenVPN use authenticated 128-bit AES, while Hamachi does not allow configuration of the security parameters and uses the default Hamachi settings.




Fig. 6. System transaction rate for various VPN approaches.

Figures 6 and 7 present the results for latency and bandwidth respectively. Latency is measured in transactions of successful request/reply messages. In the latency test, it is obvious that having the central server increases the delay between the client and server and the results degrade more quickly as additional peers are added to the system. In small systems, OpenVPN shines probably due to optimized software, though




Fig. 7. System bandwidth for various VPN approaches.

as the system grows, the system bandwidth does not. By the time 8 peers have entered into the system, both decentralized approaches perform better than the OpenVPN solution. To summarize, decentralized VPN approaches provide better scalability, which can be immediately noticed by low latency times and, as the system grows, available bandwidth.

VII. CONCLUSIONS

This paper describe all the necessary components to design and implement a decentralized VPN approach designed to facilitate ease of use while maintaining security and scalability. The methods in this paper define an abstraction for the VPN component and the features required in a structured overlay to support this style of VPN. The VPN abstraction handles local networking components and interacts with the overlay for communication, link creation, and peer discovery. Thus a structured overlay should have a DHT and potentially have NAT traversal. The paper describes an enhancement to structured overlays, autonomic relays, that enable scalable NAT traversal when NAT hole punching does not work. In addition, we presented novel ways of configuring the VPN through a Web 2.0 environment.

This approach has been used to construct real systems, such as the GroupVPN [31] and a SocialVPN [32]. The GroupVPN has been used to construct a Grid Appliance that enables the creation of distributed, decentralized, dynamic computing grids. Over the past 2 years, we have had an active grid deployed for computer architecture research, Archer [33]. Archer currently spans four universities with 500 resources, we have had 100s of users who connect seamlessly to these resources from home, school, hotels, etc. Most recently, a grid at La Jolla Institute for Allergy and Immunology went live with minimal communication with us. Researchers at the Clemson University and Purdue have opted for this approach over centralized VPNs as the basis of their future distributed compute clusters and have actively tested networks of over 1000 nodes.

In Archer, users first create an account on a web site and then request to join the Archer GroupVPN group. Once access has been granted, users can obtain configuration data that enables them to seamlessly add resources to the grid by hand-

ing the configuration file to the Grid Appliance initialization scripts. Other forms of grid environments, typically have a shared submission site, but Archer is open to all architecture researchers and monitoring usage on a shared system would require a dedicated administrator. Each user instantiates their own submission site, typically by using a virtual machine appliance, that connects to the rest of the system through the GroupVPN, enabling them to connect from anywhere without configuration. Archer allows users to easily add resources using similar mechanisms. In the case of decentralized VPNs, this would be difficult as the user would need to create a manual link to the rest of the system for each new resource. N2N may work, but the network size of Archer is larger than the recommendations made by N2N and would still require the setup of address allocation facilities. In general, all existing approaches would fail besides those with centralized components, because, at the time of this writing, all of Archer's resources are behind NATs. While centralized could be used, they would require additional dedicated resources and management and limit access if the central component went offline. Archer and Grid Appliance would not be possible without the GroupVPN architecture.

REFERENCES

- [1] J. Yonan. (2009) OpenVPN. <http://openvpn.net/>.
- [2] LogMeIn. (2009) Hamachi. <https://secure.logmein.com/products/hamachi2/>.
- [3] K. Petric. (2009, August) Wippien. <http://wippien.com/>.
- [4] G. LLC. (2009, September) Gbridge. <http://www.gbridge.com>.
- [5] A. Rezmerita, T. Morlier, V. Neri, and F. Cappello, "Private virtual cluster: Infrastructure and protocol for instant grids," in *Euro-Par '06*.
- [6] M. Tsugawa and J. Fortes, "A virtual network (vine) architecture for grid computing," *International Parallel and Distributed Processing Symposium*, vol. 0, p. 123, 2006.
- [7] X. Jiang and D. Xu, "Violin: Virtual internetworking on overlay," in *In Proc. of the 2nd Intl. Symp. on Parallel and Distributed Processing and Applications*, 2003, pp. 937–946.
- [8] A. I. Sundararaj and P. A. Dinda, "Towards virtual networks for virtual machine grid computing," in *VM'04: Proceedings of the 3rd conference on Virtual Machine Research And Technology Symposium*. USENIX Association, 2004, pp. 14–14.
- [9] G. Sliepen. (2009, September) tinc. <http://www.tinc-vpn.org/>.
- [10] L. Deri and R. Andrews, "N2N: A layer two peer-to-peer vpn," in *AIMS '08: Proceedings of the 2nd international conference on Autonomous Infrastructure, Management and Security*, 2008.
- [11] M. Krasnyansky. (2005) Universal tun/tap device driver. <http://vtun.sourceforge.net/tun>.
- [12] R. Droms, *RFC 2131 Dynamic Host Configuration Protocol*, March 1997.
- [13] W. Ginolas. (2009) P2PVPN. <http://p2pvpn.org>.
- [14] E. Exa. (2009, September) CloudVPN. <http://e-x-a.org/?view=cloudvpn>.
- [15] M. Castro, M. Costa, and A. Rowstron, "Debunking some myths about structured and unstructured overlays," in *NSDI'05: Proceedings of Symposium on Networked Systems Design & Implementation*.
- [16] A. Rowstron and P. Druschel, "Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems," in *International Conference on Distributed Systems Platforms*, November 2001.
- [17] I. Stoica and et al., "Chord: A scalable Peer-To-Peer lookup service for internet applications," in *SIGCOMM*, 2001.
- [18] G. S. Manku, M. Bawa, and P. Raghavan, "Symphony: distributed hashing in a small world," in *USITS*, 2003.
- [19] P. Maymounkov and D. Mazières, "Kademlia: A peer-to-peer information system based on the XOR metric," in *IPTPS '02*, 2002.
- [20] S. Ratnasamy, P. Francis, S. Shenker, and M. Handley, "A scalable content-addressable network," in *In Proceedings of ACM SIGCOMM*, 2001.
- [21] P. O. Boykin and et al., "A symphony conducted by brunet," <http://arxiv.org/abs/0709.4048>, 2007.
- [22] A. Ganguly, D. Wolinsky, P. Boykin, and R. Figueiredo, "Decentralized dynamic host configuration in wide-area overlays of virtual workstations," in *International Parallel and Distributed Processing Symposium*, March 2007.
- [23] I. Stoica, D. Adkins, S. Zhuang, S. Shenker, and S. Surana, "Internet indirection infrastructure," *IEEE/ACM Transactions on Networking*, 2004.
- [24] D. I. Wolinsky, Y. Liu, P. S. Juste, G. Venkatasubramanian, and R. Figueiredo, "On the design of scalable, self-configuring virtual networks," in *IEEE/ACM Supercomputing 2009*, November 2009.
- [25] D. I. Wolinsky, K. Lee, T. W. Choi, P. O. Boykin, and R. Figueiredo, "Virtual private overlays: Secure group communication in NAT-constrained environments," January 2010.
- [26] E. Rescorla and N. Modadugu. (2006, April) RFC 4347 datagram transport layer security.
- [27] J. Rosenberg, "Interactive connectivity establishment (ICE): A protocol for network address translator (NAT) traversal for offer/answer protocols," <http://tools.ietf.org/html/draft-ietf-mimusic-ice-19>, October 2008.
- [28] K. P. Gummadi, S. Saroiu, and S. D. Gribble, "King: Estimating latency between arbitrary internet end hosts," in *SIGCOMM IMW '02*.
- [29] B. Chun, D. Culler, T. Roscoe, A. Bavier, L. Peterson, M. Wavrzonik, and M. Bowman, "Planetlab: an overlay testbed for broad-coverage services," *SIGCOMM Comput. Commun. Rev.*, 2003.
- [30] J. Kleinberg, "The small-world phenomenon: an algorithm perspective," in *STOC '00: Proceedings of the thirty-second annual ACM symposium on Theory of computing*, 2000.
- [31] D. I. Wolinsky and R. Figueiredo. (2009, September) Grid appliance user interface. <http://www.grid-appliance.org>.
- [32] R. Figueiredo, P. O. B. Boykin, P. St. Juste, and D. Wolinsky, "Social vpns: Integrating overlay and social networks for seamless p2p networking."
- [33] R. Figueiredo and et al., "Archer: A community distributed computing infrastructure for computer architecture research and education," in *CollaborateCom*, November 2008.