

On the Design and Implementation of Autonomic, P2P GroupVPNs

David Isaac Wolinsky, P. Oscar Boykin, Renato Figueiredo
University of Florida

Abstract—Virtual private networks (VPNs) enable safe and simple use of network applications in distributed, insecure, and constrained environments by providing all-to-all connectivity in a secure, isolated virtual environment. Existing systems present challenges to users in small and medium groups, such as academic, business, and home environments, lacking networking, security, and operating system expertise. In this paper, we describe a VPN architecture motivated by the challenges with existing systems and their application in these environments. Our contribution is a system that provides autonomic, decentralized VPNs using common Web 2.0 group collaboration systems for configuration and management. To evaluate these claims, we provide a reference implementation, GroupVPN, and compare it to existing VPN approaches. GroupVPN is an extension of previous work known as IPOP.

I. INTRODUCTION

A Virtual Private Network (VPN) provides the illusion of a local area network (LAN) spanning a wide area network (WAN) by creating encrypted and authenticated, secure¹ communication links amongst participants. Common uses of VPNs include secure access to enterprise network resources from remote/insecure locations, connecting distributed resources from multiple sites, and establishing virtual LANs for multiplayer video games and media sharing over the Internet.

The architecture described in this paper addresses usage scenarios where VPNs are desired but complexity in deployment and management limits their applicability, such as collaborative academic environments linking individuals spanning multiple institutions, where coordinated configuration of network infrastructure across the different sites is often impractical. Another example is the small/medium business (SMB) environment, where it is often desirable to interconnect desktops and servers across distributed sites and secure traffic to enterprise networked resources without incurring the complexity or management costs of traditional VPNs. Alternatively, use of a VPN across an extended family enables sharing of media, such as family videos and pictures, though the individual sites may lack the capability of hosting a centralized service nor want computers left on all the time.

In this paper, we discuss an architecture that deals with the following VPN challenges:

- 1) **Configuration:** How will users connect to the VPN, what type of security credentials will be used, and what will the network parameters be?

- 2) **Organization:** How will peers find each other? What methods will be used for address allocation and peer discovery?
- 3) **Privacy:** Will communication in the VPN be transmitted in clear text amongst members but secure in transit or will secure links be established between source and destination?
- 4) **Connectivity:** Given a hostname for a remote node, how will the VPN discover the remote peer and what will the method used for communication?
- 5) **Peer Management:** How are new users added to the system, what if they have many resources to add as well? What happens when a user becomes malicious, how are users revoked?
- 6) **Efficiency:** How will peers communicate when they cannot form direct connections? What are efficient methods for peers to transmit broadcast / multicast IP traffic?

The work in this paper is motivated from our Archer [1] project. The primary goal for Archer is to create a dynamic and decentralized environment for computer architecture researchers to share and access voluntary compute cycles from each other. Centralized systems would limit the scope of Archer and require dedicated administration from multiple parties if the system becomes large. Decentralized VPNs typically require manual configuration of links between peers, way beyond the scope of the target users. Current P2P VPN approaches either lack scalability or proper security components to be useful for VPN approaches.

In this paper, we present a novel VPN architecture that uses an existing public overlay to self-organize a private VPN overlay using a set of P2P systems known as structured overlays and configured with a Web 2.0 group collaborative environment. We describe how the VPN can be configured from a web interface and use features of P2P systems to create an autonomic and decentralized GroupVPN. We then focus on enhancements that can be made to the VPN made available by the use of structured overlays in terms of traffic relaying and IP broadcast / multicast.

Section II introduces VPNs and various approaches to VPN overlays. Section III describes the application of core VPN features to a structured overlay. The following sections, IV and V have a detailed focus on P2P network constraints and efficient IP broadcast and multicast. Section VI presents a validation of our approach in comparison to existing approaches through quantitative analysis. We conclude in Section VII by exploring a usage scenario and how the constraints apply to our approach with existing approaches use in a real environment.

¹For the remainder of this paper, unless explicitly stated otherwise, security implies encryption and mutual authentication between peers.

II. VIRTUAL PRIVATE NETWORKS

VPNs can be divided into two components, a local or end point configuration and a network configuration. Most VPNs share a common local configuration and differ significantly in the area of network configuration. Local configuration deals with user or system interaction with the VPN, whereas network configuration tackles with the challenges of infrastructure and peer discovery, connectivity, and security. This section begins by reviewing common local configuration features and then reviews common network configurations as presented in Table I.

TABLE I
VPN CLASSIFICATIONS

Type	Description
Centralized	Clients communicate through one or more servers which are statically configured
Centralized Servers / P2P Clients	Servers provide authentication, session management, and optionally relay traffic; peers may communicate directly with each other via P2P links if NAT traversal succeeds
Decentralized Servers and Clients	No distinction between clients and servers; each member in the system authenticates directly with each other; links between members must be explicitly defined
Unstructured P2P	No distinction between clients and servers; members either know the entire network or use broadcast to discover routes between each other
Structured P2P	No distinction between clients and servers; members are usually within $O(\log N)$ hops of each other via a greedy routing algorithm; use distributed data store for discovery

A. The Basic Client VPN Configuration

Figure 1 presents an abstraction of the common features found in VPN clients: a service that communicates with the VPN system and a virtual network (VN) device for host integration. During initialization, the VPN service authenticates with the overlay² and queries it for information about the network, such as network address space, address allocations, and domain name service (DNS) servers. At which point, the VPN enables secure communication amongst participants.

A VN device allows applications to communicate transparently over the VPN. By providing mechanisms for injecting incoming packets into and retrieving outgoing packets from the networking stack, a VN device enables the use of common network APIs such as Berkeley Sockets, thereby allowing unmodified applications to work over the VPN. To do this VN devices allow the creation of a virtual network interface providing either a virtual Ethernet or IP device. The most widely available, free VN device is TAP [2]. Most operating systems support TAP either through existing features in the OS or through third party drivers, these systems include Windows, Linux, Mac OS/X, BSD, and Solaris.

VN devices can be configured manually through command-line tools or OS' APIs or dynamically by the universally supported dynamic host configuration process (DHCP) [3]. When a VN device obtains an IP address, it triggers the OS

to add a new rule to the routing table directing all packets sent to the VPN subnet to the VN device. Packets read from the TAP device are encrypted and sent to the overlay via the VPN client. The overlay delivers the packet to another client or a server with a VN stack enabled. Received packets are decrypted, verified for authenticity, and then written to the VN device. In most cases, the IP layer header remains unchanged, while VPN configuration determines how the Ethernet header is handled.

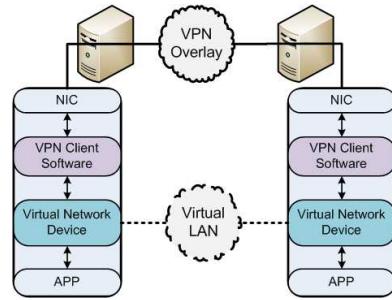


Fig. 1. A typical VPN client. A VN device makes application interaction with the VPN transparent by integrating with the OS network stack.

B. Centralized VPN Systems

OpenVPN is an open and well-documented platform for deploying centralized VPNs. In this paper, it is used as the basis for understanding centralized VPNs as it represents features common to most centralized VPNs.

In centralized VPN systems, clients forward all VPN related packets to the server. Client responsibilities are limited to configuring the VN device and authenticating with the VPN server; whereas the servers are responsible for authentication and routing between clients. Likewise, broadcast and multicast packets also must pass through the central server.

Centralized systems do not prohibit multiple servers, though each server must be manually configured to interact with remote servers. Clients randomly select a server from a list of known servers, implementing a simple load balance. Once connected, the servers provide the client an IP address in the VPN address space. Thus two peers on different servers will have three hops between each other, two from routing over the two servers and one for the forwarding to the destination.

All inter-client communication flows through a central server. By default, a client encrypts a packet and sends it to the server. Upon receiving the packet, the server decrypts it, determines where to relay it, encrypts it, and then sends the packet to its destination. This model allows a server to eavesdrop on communication. While a second layer of encryption is possible through a shared secret, it requires out-of-band communication and increases the computing overhead on communication since the packet will now need to be encrypted and authenticated twice at both the VPN source and destination sites.

C. Centralized P2P VPN Systems

Hamachi [4] is the first well-known centralized VPN that used the ambiguous moniker "P2P VPN". In reality, these

²An overlay in this context refers to the topology formed by VPN nodes including central servers, other VPN clients, or relays.

systems are better classified as centralized VPN servers with P2P links. Similar VPNs include Wippien [5], Gbridge [6], PVC [7], and P2PVPN³ [8]. The P2P in these systems is limited to direct connectivity between clients orchestrated through a central server: in Wippien it is a chat server, while P2PVPN uses a BitTorrent tracker. If NAT traversal or firewalls prevent direct connectivity, the central server can act as a relay. Each approach uses their own security protocols with most using a server to verify the authenticity and setup secure connections between clients. Thus the peers must trust the central authentication server, though communication between peers, or EtE communication, is secure. In regards to the P2PVPN, long term goals involve the creation of an unstructured network, which would provide a method of decentralized organization. More security minded individuals may feel uncomfortable using Hamachi and GBridge as they do not support independent systems, thus requiring use of their private resources.

D. Decentralized VPN Systems

Some examples of systems that assist in distributing load in VPN systems are tinc [9], CloudVPN [10], ViNe [11], VNET [12], and Violin [13]. These systems are not autonomic and require explicit configuration of links between resources. This means that, like OpenVPN, these systems can suffer VPN outages when nodes go offline, thus administrators must maintain the VPN connection table. Unlike OpenVPN, these approaches typically do not require all-to-all direct connectivity for all-to-all communication. These systems support NAT traversal through out-of-band mechanisms or routing through relays. Current decentralized VPNs only secure PtP communication, leaving EtE communication in the clear text when sent via the overlay.

E. Unstructured P2P VPN Systems

Unlike centralized and decentralized systems, P2P environments require the user to connect to the overlay, which then automatically configures links. The simplest form of overlays are unstructured, where peers form random connections with each other and use broadcast and stochastic techniques to find information and other peers, though due to its unstructured nature, the system cannot guarantee distance and routability between peers. The only example of an unstructured VPN is N2N [14]. In N2N, peers first connect to a super node and then, to find another peer, they broadcast discovery messages to the entire overlay. In the case that peers cannot form direct connection, peers can route to each other over the N2N overlay. In the realm of VPNs, all client VPNs are also servers performing authentication though neither approach deals with decentralized address allocation. N2N does not address the challenges of IP address allocation and management, relegating that to the user to handle.

³Due to the similarities between the name P2PVPN and focus of this paper, “P2PVPN” refers to [8] and “P2P VPN” to the use of P2P in VPNs.

III. STRUCTURED P2P VPN SYSTEMS

This section describes the basic construction of a Structured P2P VPN, including background on structured overlays, address allocation and discovery, security, and connectivity.

A. Background

An alternative to dealing with scalability concerns in unstructured systems, structured overlays provide distributed look up services with guaranteed search time with a lower bound of $O(\log N)$, in contrast to unstructured systems, which rely on global knowledge/broadcasts, or stochastic techniques such as random walks [15]. Some examples of structured systems are Pastry [16], Chord [17], Symphony [18], Kademia [19], CAN [20], and Brunet [21]. In general, structured systems make these guarantees by self-organizing into a well-defined topology, such as a 2D ring or a hypercube, by randomly generated, though uniformly distributed, node identifiers. A key component of most structured overlays is support for decentralized storage/retrieval of information by mapping keys to specific node IDs in an overlay called a distributed hash table (DHT). At a minimum, the data is stored at the node ID either smaller or larger to the data’s node ID and for fault tolerance the data can be stored at other nodes. DHTs can be used by peers of systems to coordinate allocation and discovery of resources, making them attractive for self-configuration in decentralized collaborative environments.

B. Organization

In the context of VPNs, structured overlays can handle organization of the network space, address allocation and discovery, decentrally through the use of a DHT, such systems have been proposed in [22], [23]. Membership in the VPN includes a matching membership in the structured overlay, thus all VPN peers have a node ID. To address the challenges of having multiple VPNs in the same overlay, each VPN has its own namespace, reducing the likelihood of overlap. To enable scalable and decentralized address allocation and discovery, peers store mappings of IP address to node ID into the DHT, typically of the form $\text{hash}(\text{namespace} + \text{IP}) = \text{nodeID}$. Thus a peer attempting to allocate an address will insert this key, value pair into the overlay. The first peer to do this will be the owner of the IP address allocation. Therefore the DHT must either support atomic writes into the DHT and prevent future writes or provide creation times of key, value pairs during queries.

Mechanisms to self-configure the IP address and network parameters of the local system can be provided by DHCP, manually configuring the IP address, or the VPN hooking into OS APIs. Address discovery is initiated when an outgoing packet for a remote peer arrives at the VPN software. At which point, the VPN will query the DHT with the IP address to obtain the owner’s node ID and forward the packet to the destination. Discussion on both these topics is further covered in our previous work [24].

C. Privacy for the VPN

Structured overlays are difficult to secure. Malicious users can pollute the DHT, send bogus messages, and even prevent the overlay from functioning. When using a structured overlay for a VPN, tamper-resistance and reliable routing mechanisms are ideal, but they only reduce the ability of a malicious user to affect the overlay, an expectation of a VPN is that no members are malicious and malicious users have their access revoked. As such a structured overlay requires both the point-to-point (PtP), communication between connected peers, and end-to-end (EtE), communication between peers across the overlay, security. Each component provides security uniquely from each other. PtP security ensures that malicious third parties cannot insert material into the system, while EtE ensures message security between the communicating peers. Enabling just PtP security would limit the security to the type of VPN security found in other VPNs, all members must be trusted, whereas EtE enables message security even from overlay members.

In our previous work [25], we described mechanisms that enable the construction of a private overlay providing both EtE and PtP security. In this system, EtE and PtP communication has been abstracted into a common sender layer; while the security component has been placed into a filter, enabling it to be placed anywhere in the sending stack and thus securing either EtE or PtP communication without additional complexity. Existing security filter supports DTLS [26], as using DTLS does not require reliable communication across the overlay or between peers.

D. Connecting to the VPN

As of 2010, the majority of the Internet is connected via Internet Protocol (IP) version 4. A limitation in this protocol is that there are only 2^{32} addresses (approximately 4 billion) available. With the Earth's population of over 8 billion and each individual potentially having multiple devices that have Internet connectivity, the IPv4 limitation is becoming more and more apparent. To address this issue there have been two approaches: 1) the use of network address translation (NAT) enabling many devices to share a single IP address though limiting connection initiation, and 2) IPv6 which supports 2^{128} addresses but comes with no guarantees of the demise of NATs or firewalls that only allow outbound connections.

NATs and firewalls are impediments to P2P systems as they limit all-to-all connectivity. There has been much research to address these challenges, the results come in the form of two types of NAT traversal: hole punching and relaying. When performing hole punching [27], peers exchanged IP and port information through a third-party proxy and simultaneously attempt to create "holes" in their NAT / firewall devices allowing direct TCP or UDP communication with a peer behind another NAT / firewall. This works by confusing the NAT into believing that the peer behind it initiated the connection. TCP NAT hole punching, though, does not work on systems that use stateful firewalls. Relaying, on the other hand, always works as peers behind NATs communicate through a third peer

with whom they have direct connectivity, though with the side-effect of additional latency and the cost incurred by the relay of supporting the bidirectional communication.

NATs are not the only consideration for the targeted environments, a reasonable expectation is that many users may be unaware of their local IP nor understand concepts such as ports. As such, it is unsafe to assume these peers can properly act as a bootstrapping peer in a small P2P system. So even when peers have public IP addresses, the owner may be unaware of how or unwilling to properly utilize it.

These considerations lead into the common bootstrapping problem for P2P systems. Creating and connecting to a private overlay can be difficult when there are no dedicated resources with public addresses, even if the system has many active members behind NATs and firewalls. This is because peers behind NATs and firewalls may be unable to handle incoming connections without the assistance of a third-party. The same work [25] describes how arbitrary public overlays can act as a middle ground to bootstrapping private overlays, using mechanisms unique to these environments to bootstrap into the private overlay. As mentioned in the introduction, peers can discover each other using public overlays to connect to a private overlay. Examples of potential, large public overlays include Gnutella, Skype, and Kademia. The key challenge to using public overlays is that they must allow peers to send arbitrary data to each other via the overlay, otherwise two peers behind NATs will be unable to coordinate NAT traversal with each other. Of course, this mechanism does not preclude the creation of a private overlay without the use of a public bootstrapping mechanism.

In practical implementation, peers first join a public structured overlay using the public overlay's DHT to find other peers. The private overlay node can then either attempt to connect directly with other private overlay nodes using TCP or UDP or use the public overlay as another transport mechanism. We have termed this approach to creating private overlays, virtual private overlays. In all cases, the communication is secured via the EtE and PtP security filters. This approach enables the construction of an overlay consisting entirely of peers behind NATs and firewalls, who would otherwise be unable to form their own overlay.

E. Managing and Configuring the VPN

A Group based Web 2.0 environment enables low overhead configuration of collaborative environments. The roles in a group environment can be divided into administrators and users. Users have the ability to join and create groups; whereas administrators define network parameters, accept or deny join requests, remove users, and promote other users to administrators. By applying this to a VPN, the group environment provides a wrapper around a public key infrastructure (PKI), where the administrators of the group act as the certificate authority (CA) and the members have the ability to obtain signed certificates. Elaborating further, when a user joins a group, the administrator can enable automatic signing of certificates or require prior review; and when peers have

overstayed their welcome, an administrator can revoke their certificate by removing them from the group. As described in [25], revocation can be stored as a CRL (certificate revocation list) on the group site or distributed through broadcast and DHT on the overlay. Though for these forms of systems a user revocation list (URL) as opposed to a CRL simplifies revocation in the case where a malicious user has automatically signed large amounts of certificates.

In the case of the GroupVPN, administrators of a group configure the VPN address range, namespace, and security. Options are also available to specify the reuse of an existing public overlay or a list of user managed nodes. When a user has been accepted into the group, they are able to download VPN configuration data, which can be seamlessly added to the VPN by running the VPN configuration process. In addition to IP address range, namespace, and security options, the configuration data also contains the group website's address and a shared secret. The shared secret uniquely identifies the user, so that the website can automatically sign the certificate or enqueue it so the administrator can manually sign it. When making a certificate request, the user sends over HTTPS a public key and their shared secret, the website creates and signs a certificate request based upon the public key and the user's relevant information ensuring that users cannot trick the website into signing malicious data. Upon receiving the signed certificate, peers are able to join the private overlay and VPN enabling secure communication amongst the VPN peers.

There are many ways of implementing and hosting the web site. For example, Google offers free hosting of Python web applications through Google Apps, but this requires that the user owns a domain. Alternatively, the user could host the group site on a public VN, peers interacting with the GroupVPN would need to connect with the public VN in order to create an account, get the configuration data, and to sign their certificate at which point they could disconnect from it. This does not preclude the use of other social mediums nor a central site dedicated to the formation of many GroupVPNs. Many GroupVPNs can share a single site, so long as the group members trust the site to host the CA private key.

IV. EFFICIENTLY DEALING WITH NATS AND FIREWALLS

Centralized and decentralized VPNs do not suffer from connectivity limitations due to NATs and firewalls as all traffic passes through the central server or managed links. In the centralized P2P approaches, peers that cannot form direct links can potentially route through a dedicated relay system, but if not, the peers will be unable to communicate. Whereas in decentralized approaches, including P2P systems, peers can communicate through the overlay when traversal fails. As the overlay grows in size, especially in structured overlays, so does the hop count between peers, increasing latency and reducing bandwidth, greatly reducing the value of routing through an overlay.

A. An Autonomic Relay Infrastructure

P2P systems can address this connectivity issue through a distributed, autonomic relaying system, whereby peers form

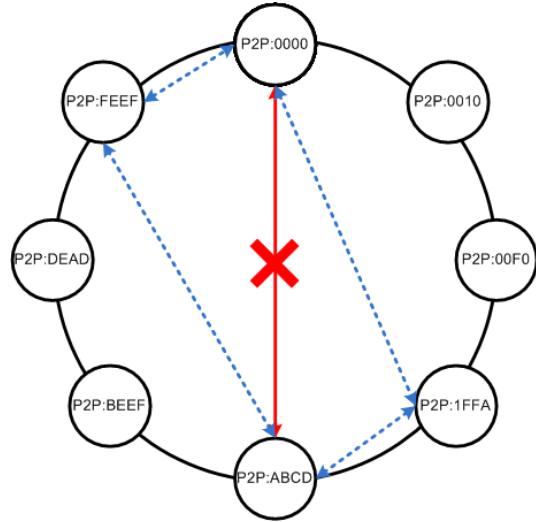


Fig. 2. Creating relays when direct connectivity is not possible. Two members, 0000 and ABCD, are unable to form direct connectivity, thus they exchange neighbor information through the overlay and connect to one of each other's neighbors, creating an overlap. The overlap then becomes a relay path (represented by dashed lines) for two-hop communication.

overlapping connections with each other's peers, creating two-hop links. When two nodes discover each other via the overlay and attempt to become connected but fail during NAT and firewall traversal, they exchange peer lists through the overlay. Upon receiving this list, the two peers identify the overlap in the neighbor sets to form a two-hop connection. When peers are not located near each other in the overlay, most likely they will have no overlap. This can be addressed by having peers connect to each other's neighbor set proactively creating overlap, as represented in Figure 2.

With the set of neighbors, the peers can include additional data such as latency and connection lifetime to each of the neighbors. When creating overlap or overlap changes, peers review these metrics to decide which subset of the overlap to use as a proxy leaving the remaining peers as reserves.

B. Usefulness of Relays

To determine the value in two-hop overlays as opposed to using the overlay in restrictive NAT and firewall scenarios, this experiment uses a network modeler that reuses the code base of the GroupVPN to faithfully implement routing in an overlay as well as latency between peers. Pair-wise latency in this experiment is set using the MIT distribution of the King Data Set [28], consisting of all-to-all latency for 1,740 distributed Internet peers. The modeler randomly assigns each peer in the modeler a matching physical location based upon the data set, thus for simulations of over 1,740, peers may be co-located. The modeler measured average overlay, 2-hop, and 1-hop latency between peers over 100 executions. Overlay latency represents the time to route a packet via the overlay. 2-hop latency is based upon the low-latency formed 2-hop connections, where peers route over the overlap with the lowest latency. 1-hop latency is based upon peers forming a direct connection with each other. Only peers that are not directly connected (i.e., have two hops or more and are not

located in the same physical location) are analyzed, as these situations would only benefit under triangular inequalities, which are not a consideration in this work.

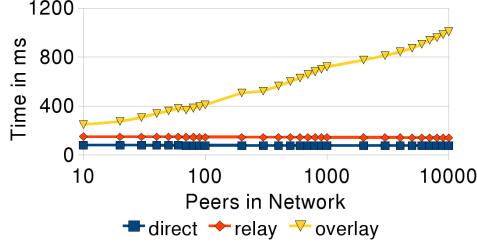


Fig. 3. A comparison of the average all-to-all overlay routing, two-hop relay, and direct connection latency in a ring based structured overlay.

Results are presented in Figure 3. Immediately the two-hop routing shows improvement. While the overlay grows logarithmically, the results for direct and relay remain nearly identical. Clearly when timing, which can affect bandwidth, is critical, direct paths are most desired, though when they are not available, this experiment clearly shows that two-hop paths, or relays, are significantly better than routing the packet through an overlay.

In addition, the relays functionality was verified in GroupVPN using a real system consisting of PlanetLab as the public overlay and the Archer as the VPN environment. PlanetLab [29] provides a set of over 500 distributed computing nodes all with public IP addresses. Archer provides grid computing to computer architecture researchers and consists of over 500 nodes located at various academic sites. To ensure that peers under study formed two-hop connections, a firewall was instantiated preventing direct communication between the two. In tests, it was revealed that the overlay was always able to self-configure a relay; the bandwidth and latency averages and deviations were $2245 \text{ Kbit/s} \pm 1080$ and $58.1 \text{ ms} \pm 35.5$, respectively.

V. EFFICIENT BROADCAST AND MULTICAST

Within the past few years, self-configuring protocols like zeroconf and Bonjour have enabled the use of complex applications transparently through self-configuring; these include music sharing applications, like iTunes, chat programs, like Pidgin, and even web sites and data shares. The underlying enabling technology is multicast DNS and service discovery. Multicast has also been used in other applications to enable efficient streaming of media, such as music and videos. These types of applications and enabling are ideal for small and medium groups.

Many VPN technologies either do not support multicast and broadcast or route the packets to all peers or through a central server, which may be reasonable for relatively small groups and for discovery, but they do not scale well for media streaming or large groups. In the original set of structured P2P virtual networks, broadcast and multicast used the DHT to keep a list of all the active peers in the VPN. When a peer wanted to send out a broadcast or multicast packet, it would query this list and then send out unicast packets to each

peer either directly or via the overlay. Furthermore, as more peers join the system, so does the count of key, value pairs in the DHT. If the peer wishes to send out many multicast or broadcast messages after another, it would either need to query the DHT every time or maintain a cache. In the case of a cache, new peers joining the overlay would need to notify the entire VPN when they joined so that they could be added to caches. The simplistic approach of querying each time does not scale well, while the other approach involving caches can be quite complicated to implement.

A. Overlay Broadcast

With the use of virtual private overlays, broadcast and multicast can be implemented stateless as they can be routed using overlay broadcast. An example of a scalable ring-based structured P2P overlay broadcast algorithm was presented in [25] to broadcast certificate revocation within a private overlay. The algorithm is called bounded-broadcast, because it is capable of broadcasting to a subset of the ring or the entire overlay. Bounded broadcast uses the following recursive algorithm: Begin with node x triggering a broadcast message over the region $[x, y]$. x has F connections to nodes in the range $[x, y]$. Denoting the i^{th} such neighbor as b_i , the node x sends a bounded broadcast over a sub-range, $[b_i, b_{i+1})$, to b_i , except the final neighbor. Differently stated, b_i is in charge of bounded-broadcasting in the sub-range $[b_i, b_{i+1})$. If there is no connection to a node in the sub-range, the recursion has ended. The final neighbor, (b_F), is responsible for continuing the bounded broadcast from $[b_F, y]$. When a node receives a message to a range that contains its own address the message is delivered to that node and then routed to others in that range. Figure 4 shows how this bounded broadcast forms a local tree recursively. The time and bandwidth for a bounded broadcast is $O(\log^2 N)$ and $N * \text{packetsize}$ as shown in [30].

To perform a broadcast on the entire overlay, a peer performs the bounded-broadcast starting from its node ID with the end address being the node ID immediately preceding its own in the address space. Though in VPN situations, many peers may already have connections to most if not all of their VPN peers, thus the broadcast algorithm has been modified to route to only connections created for structured overlay purposes and not explicitly only for VPN purposes. Otherwise in many cases, this algorithm will degenerate into one similar to the unicast approach.

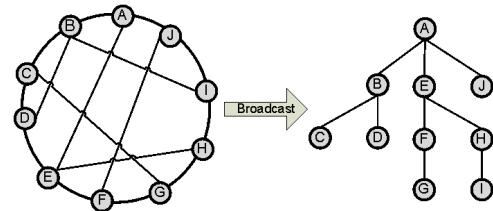


Fig. 4. Bounded-Broadcast performing a complete overlay broadcast

Another consideration, left unexplored, is whether or not to create separate overlays for each multicast group. In IP multicast, peers register group membership using IGMP (Internet

Group Management Protocol), thus a VPN can be configured to parse these packets and join a unique private overlay for that multicast group. In general, the cost of joining an additional overlay is minimal, though, when used for zeroconf like applications, most peers will join the group and more importantly the traffic is minimal anyway.

B. IP Broadcast Evaluation

The broadcast validation was done using the same modeler used in the relay evaluation. This evaluation compares broadcasting to a simple unicast mechanism that assumes the peer has already queried the DHT. Both models are evaluated in a private overlay varying from 10 members up to 10,000. It is important to note that the network modeler does not consider host overheads including bandwidth limitations nor simultaneous packet sending. The modeler measures network bandwidth, sending peer bandwidth, and the time to complete the broadcast.

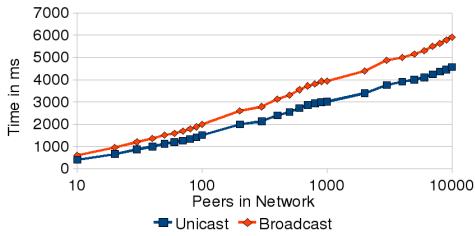


Fig. 5. Time to complete an IP Broadcast

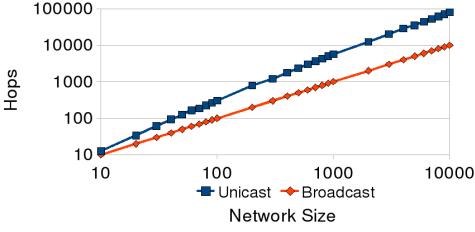


Fig. 6. Number of peers the broadcast packet traversed

Comparing the timing results, Figure 5, to the number of hops, Figure 6, indicates that the two are related by a factor of approximately $\log(N)$. This experiment shows that the broadcast performs reasonably similar to unicast, which when combined with the fact that broadcast is stateless and does not put undue stress on the sending site as the network size increases makes it clear that this method is preferred for IP broadcasting and multicasting.

VI. EVALUATION OF VPN MODELS

This experiment explores bandwidth and latency in a distributed VPN system to motivate the usage of P2P links in a VPN. The VPNs used include our GroupVPN, OpenVPN, and Hamachi. OpenVPN represents a typical centralized VPN, while Hamachi represents a well-tuned P2P-link VPN. The evaluation was performed on Amazon EC2 using small instance sized Ubuntu i386 instances to create various sized networks ranging from 1 to 32. OpenVPN uses an additional node as the central server and Hamachi has an upper bound

of 16 due to limitations in the Linux version at the time of this evaluation. To perform bandwidth tests, the instances are booted and query an NFS for the list virtual IP addresses, peers are ordered such that half the peers are act as clients and the other half the peers creating a 1 to 1 mapping between all sets. Latency and bandwidth tests are performed using netperf's request-reply and streaming tests respectively. Prior to the start of the tests, peers have no knowledge of each other, except the virtual IP addresses, thus connection startup costs are included in the test. Tests are run for 10 minutes diluting the connection initiation overhead but represent an example of real usage. Results from the clients are polled at all locations and averaged together, though the OpenVPN server is measured separately. GroupVPN and OpenVPN use authenticated 128-bit AES, while Hamachi does not allow configuration of the security parameters and uses the default Hamachi settings, 256-bit AES.

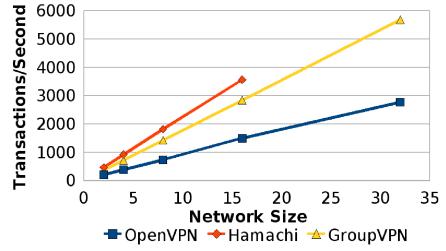


Fig. 7. System transaction rate for various VPN approaches.

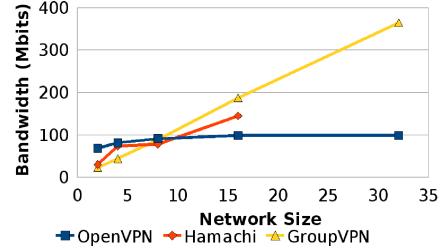


Fig. 8. System bandwidth for various VPN approaches.

Figures 7 and 8 present the results for latency and bandwidth respectively. Latency is measured in transactions of successful request/reply messages. In the latency test, it is obvious that having the central server increases the delay between the client and server and the results degrade more quickly as additional peers are added to the system. In small systems, OpenVPN shines probably due to optimized software, though as the system grows, the system bandwidth does not. By the time 8 peers have entered into the system, both decentralized approaches perform better than the OpenVPN solution. To summarize, decentralized VPN approaches provide better scalability, which can be immediately noticed by low latency times and, as the system grows, available bandwidth.

VII. CONCLUSIONS

This paper describes a novel approach to VPNs utilizing structured overlays to deal with organization, public overlays for connectivity, private overlays for security, and collaborative

environments for configuration and management. The paper furthers previous work using virtual private overlays to enhance key VPN features, relays and IP broadcast and multicast.

Without the functionality of GroupVPN projects like Archer [1], would be impractical. Archer consists of over 500 resources from 5 different universities, including University of Florida, Florida State University, Northeastern University, University of Minnesota, and University of Texas. In the past year, since Archer came online, over 100 unique users have contributed and taken advantage of the voluntary computing cycles. A new user to the system begins by creating an account at Grid-Appliance.org and requesting membership in the Archer GroupVPN group. Once access has been granted, users can obtain configuration data used by the Grid Appliance initialization scripts to seamlessly add resources to the grid. This method allows independent submission sites, unlike most grid systems that have a shared submission site, which require dedicated administrators. Most users connect to the system using a pre-configured virtual machine appliance, so that they do not need to be experts in grid systems to take advantage of Archer. Enabling this using decentralized VPNs would be difficult as the user would need to create manual links to the rest of the system for each new resource. N2N may work, but the network size of Archer is larger than the recommendations made by N2N and would still require the setup of address allocation facilities. In general, all existing approaches would fail besides those with centralized components, because, at the time of this writing, all of Archer's resources are behind NATs. Even though centralized could be used, it would require additional dedicated resources and management, limiting access if the central component went offline.

Much of the work in this paper has contributed to the framework for SocialVPN [31]. Instead of a group environment for peer management, SocialVPN, like Wippien, uses an online social networking backend, such as Facebook, MySpace, or Google Talk, to discover VPN partners. Unlike Wippien, which will then attempt to negotiate links through the social network, SocialVPN decentralizes itself and uses a public overlay to discover the peers different end points, establish direct and relayed links, and form private overlays for broadcast and multicast messages. After registering with each other through a social network, peers are able to discover and connect with each other through the overlay.

The GroupVPN has been used as the virtual network for the Grid Appliance, which is the basis of Archer and, in general, enables the creation of distributed, decentralized, collaborative environments for computing grids. Recently, a grid at La Jolla Institute for Allergy and Immunology went live using GroupVPN and Grid Appliance without receiving any technical support from us. Researchers at Clemson University and Purdue have opted for this approach over centralized VPNs as the basis of their future distributed compute clusters and have actively tested networks of over 700 nodes.

REFERENCES

- [1] R. Figueiredo and et al., "Archer: A community distributed computing infrastructure for computer architecture research and education," in *CollaborateCom*, November 2008.
- [2] M. Krasnyansky. (2005) Universal tun/tap device driver. <http://vtun.sourceforge.net/tun>.
- [3] R. Droms, *RFC 2131 Dynamic Host Configuration Protocol*, March 1997.
- [4] LogMeIn. (2009) Hamachi. <https://secure.logmein.com/products/hamachi2/>.
- [5] K. Petric. (2009, August) Wippien. <http://wippien.com/>.
- [6] G. LLC. (2009, September) Gbridge. <http://www.gbridge.com>.
- [7] A. Rezmerita, T. Morlier, V. Neri, and F. Cappello, "Private virtual cluster: Infrastructure and protocol for instant grids," in *Euro-Par '06*.
- [8] W. Ginolas. (2009) P2PVPN. <http://p2pvpn.org>.
- [9] G. Sliepen. (2009, September) tinc. <http://www.tinc-vpn.org/>.
- [10] E. Exa. (2009, September) CloudVPN. <http://e-x-a.org/?view=cloudvpn>.
- [11] M. Tsugawa and J. Fortes, "A virtual network (vine) architecture for grid computing," in *International Parallel and Distributed Processing Symposium*, vol. 0, p. 123, 2006.
- [12] A. I. Sundararaj and P. A. Dinda, "Towards virtual networks for virtual machine grid computing," in *VM'04: Proceedings of the 3rd conference on Virtual Machine Research And Technology Symposium*. USENIX Association, 2004, pp. 14–14.
- [13] X. Jiang and D. Xu, "Violin: Virtual internetworking on overlay," in *In Proc. of the 2nd Intl. Symp. on Parallel and Distributed Processing and Applications*, 2003, pp. 937–946.
- [14] L. Deri and R. Andrews, "N2N: A layer two peer-to-peer vpn," in *AIMS '08: Proceedings of the 2nd international conference on Autonomous Infrastructure, Management and Security*, 2008.
- [15] M. Castro, M. Costa, and A. Rowstron, "Debunking some myths about structured and unstructured overlays," in *NSDI'05: Proceedings of Symposium on Networked Systems Design & Implementation*.
- [16] A. Rowstron and P. Druschel, "Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems," in *International Conference on Distributed Systems Platforms*, November 2001.
- [17] I. Stoica and et al., "Chord: A scalable Peer-To-Peer lookup service for internet applications," in *SIGCOMM*, 2001.
- [18] G. S. Manku, M. Bawa, and P. Raghavan, "Symphony: distributed hashing in a small world," in *USITS*, 2003.
- [19] P. Maymounkov and D. Mazières, "Kademlia: A peer-to-peer information system based on the XOR metric," in *IPTPS '02*, 2002.
- [20] S. Ratnasamy, P. Francis, S. Shenker, and M. Handley, "A scalable content-addressable network," in *In Proceedings of ACM SIGCOMM*, 2001.
- [21] P. O. Boykin and et al., "A symphony conducted by brunet," <http://arxiv.org/abs/0709.4048>, 2007.
- [22] A. Ganguly, D. Wolinsky, P. Boykin, and R. Figueiredo, "Decentralized dynamic host configuration in wide-area overlays of virtual workstations," in *International Parallel and Distributed Processing Symposium*, March 2007.
- [23] I. Stoica, D. Adkins, S. Zhuang, S. Shenker, and S. Surana, "Internet indirection infrastructure," *IEEE/ACM Transactions on Networking*, 2004.
- [24] D. I. Wolinsky, Y. Liu, P. S. Juste, G. Venkatasubramanian, and R. Figueiredo, "On the design of scalable, self-configuring virtual networks," in *IEEE/ACM Supercomputing 2009*, November 2009.
- [25] D. I. Wolinsky, K. Lee, T. W. Choi, P. O. Boykin, and R. Figueiredo, "Virtual private overlays: Secure group communication in NAT-constrained environments," January 2010.
- [26] E. Rescorla and N. Modadugu. (2006, April) RFC 4347 datagram transport layer security.
- [27] J. Rosenberg, "Interactive connectivity establishment (ICE): A protocol for network address translator (NAT) traversal for offer/answer protocols," <http://tools.ietf.org/html/draft-ietf-mmusic-ice-19>, October 2008.
- [28] K. P. Gummadi, S. Saroiu, and S. D. Gribble, "King: Estimating latency between arbitrary internet end hosts," in *SIGCOMM IMW '02*.
- [29] B. Chun, D. Culler, T. Roscoe, A. Bavier, L. Peterson, M. Wavrzonik, and M. Bowman, "Planetlab: an overlay testbed for broad-coverage services," *SIGCOMM Comput. Commun. Rev.*, 2003.
- [30] J. Kleinberg, "The small-world phenomenon: an algorithm perspective," in *STOC '00: Proceedings of the thirty-second annual ACM symposium on Theory of computing*, 2000.
- [31] R. Figueiredo, P. O. B. Boykin, P. St. Juste, and D. Wolinsky, "Social vpns: Integrating overlay and social networks for seamless p2p networking."